

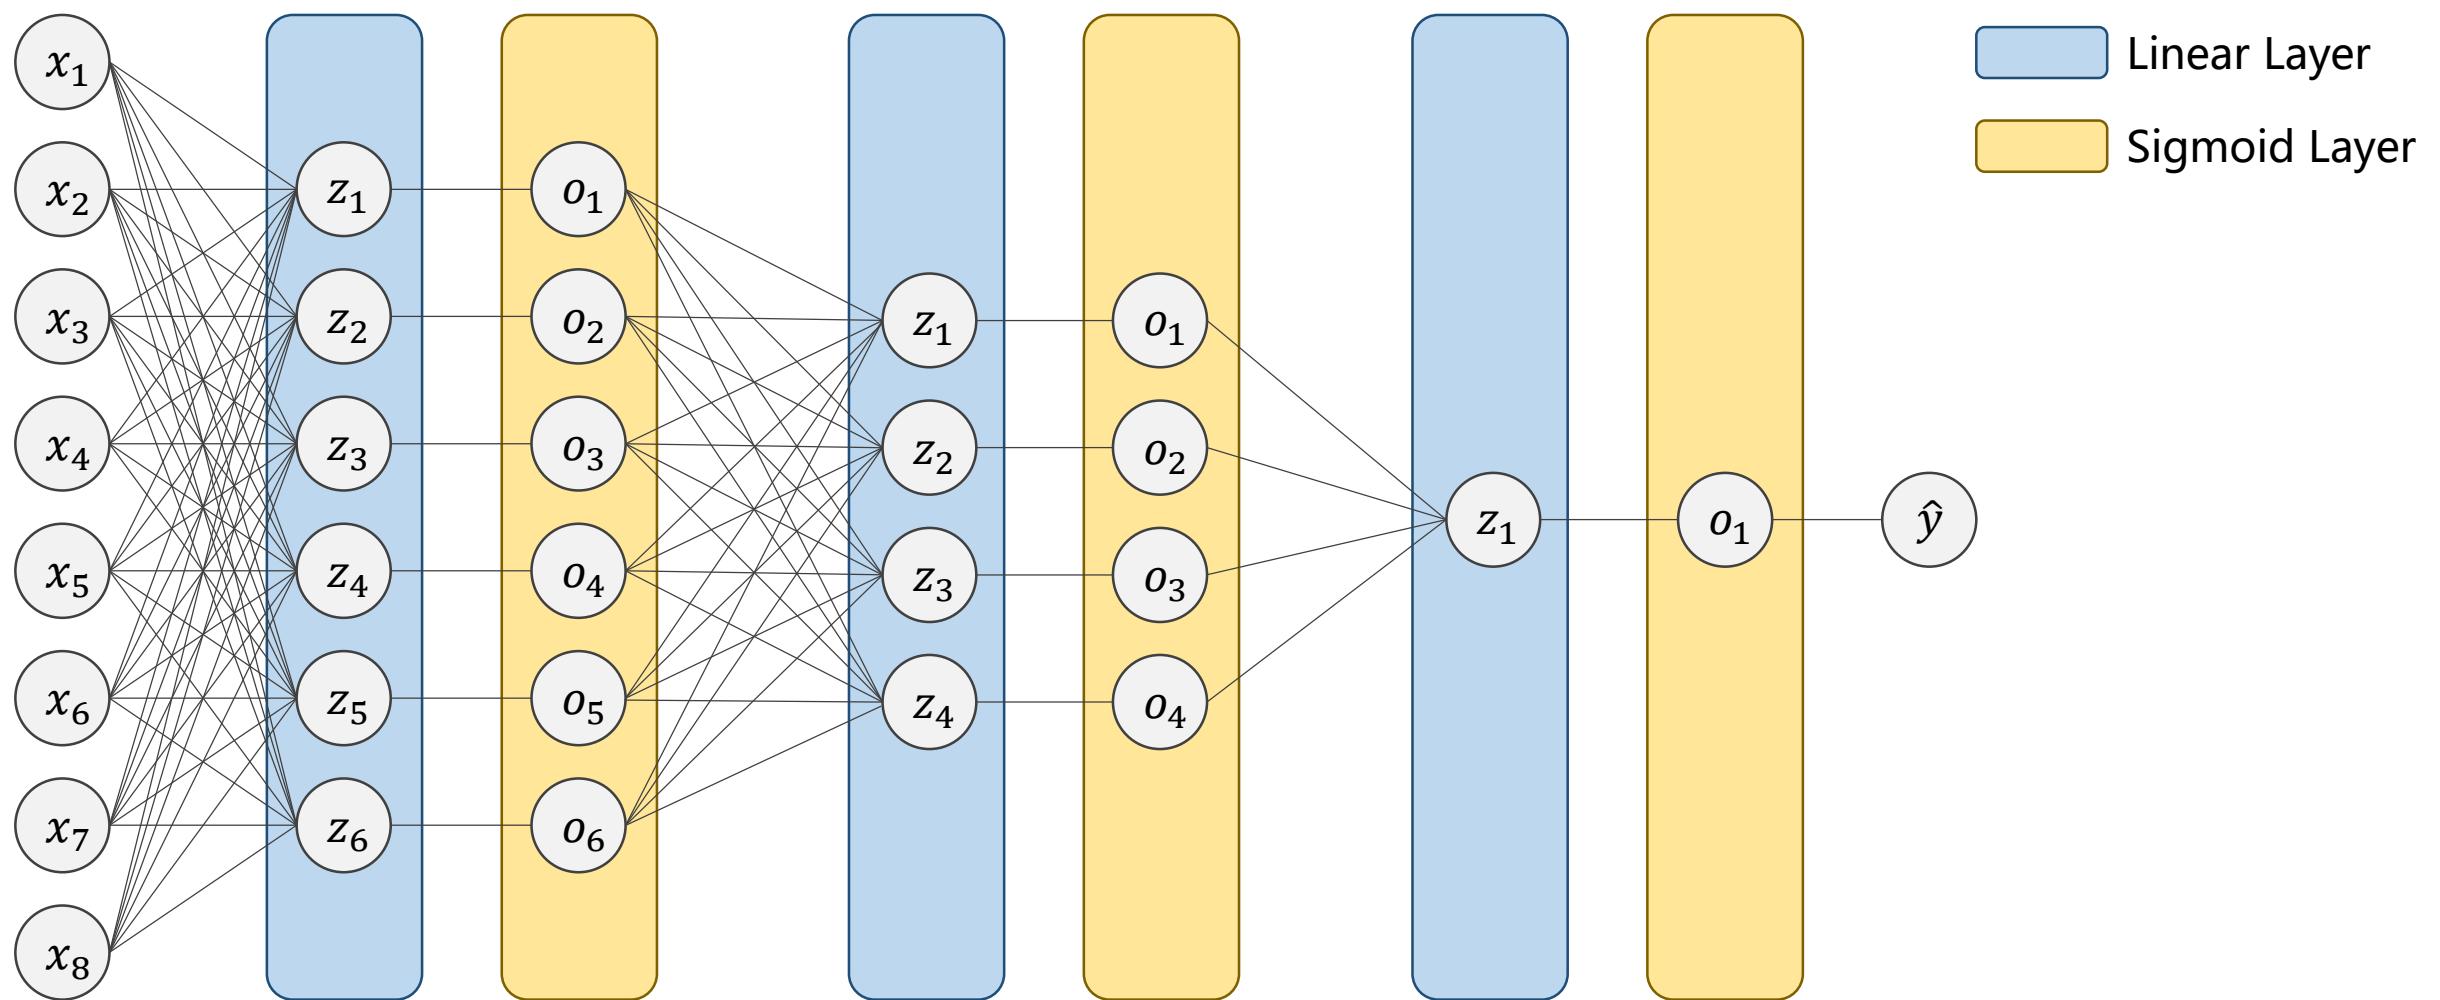


# PyTorch Tutorial

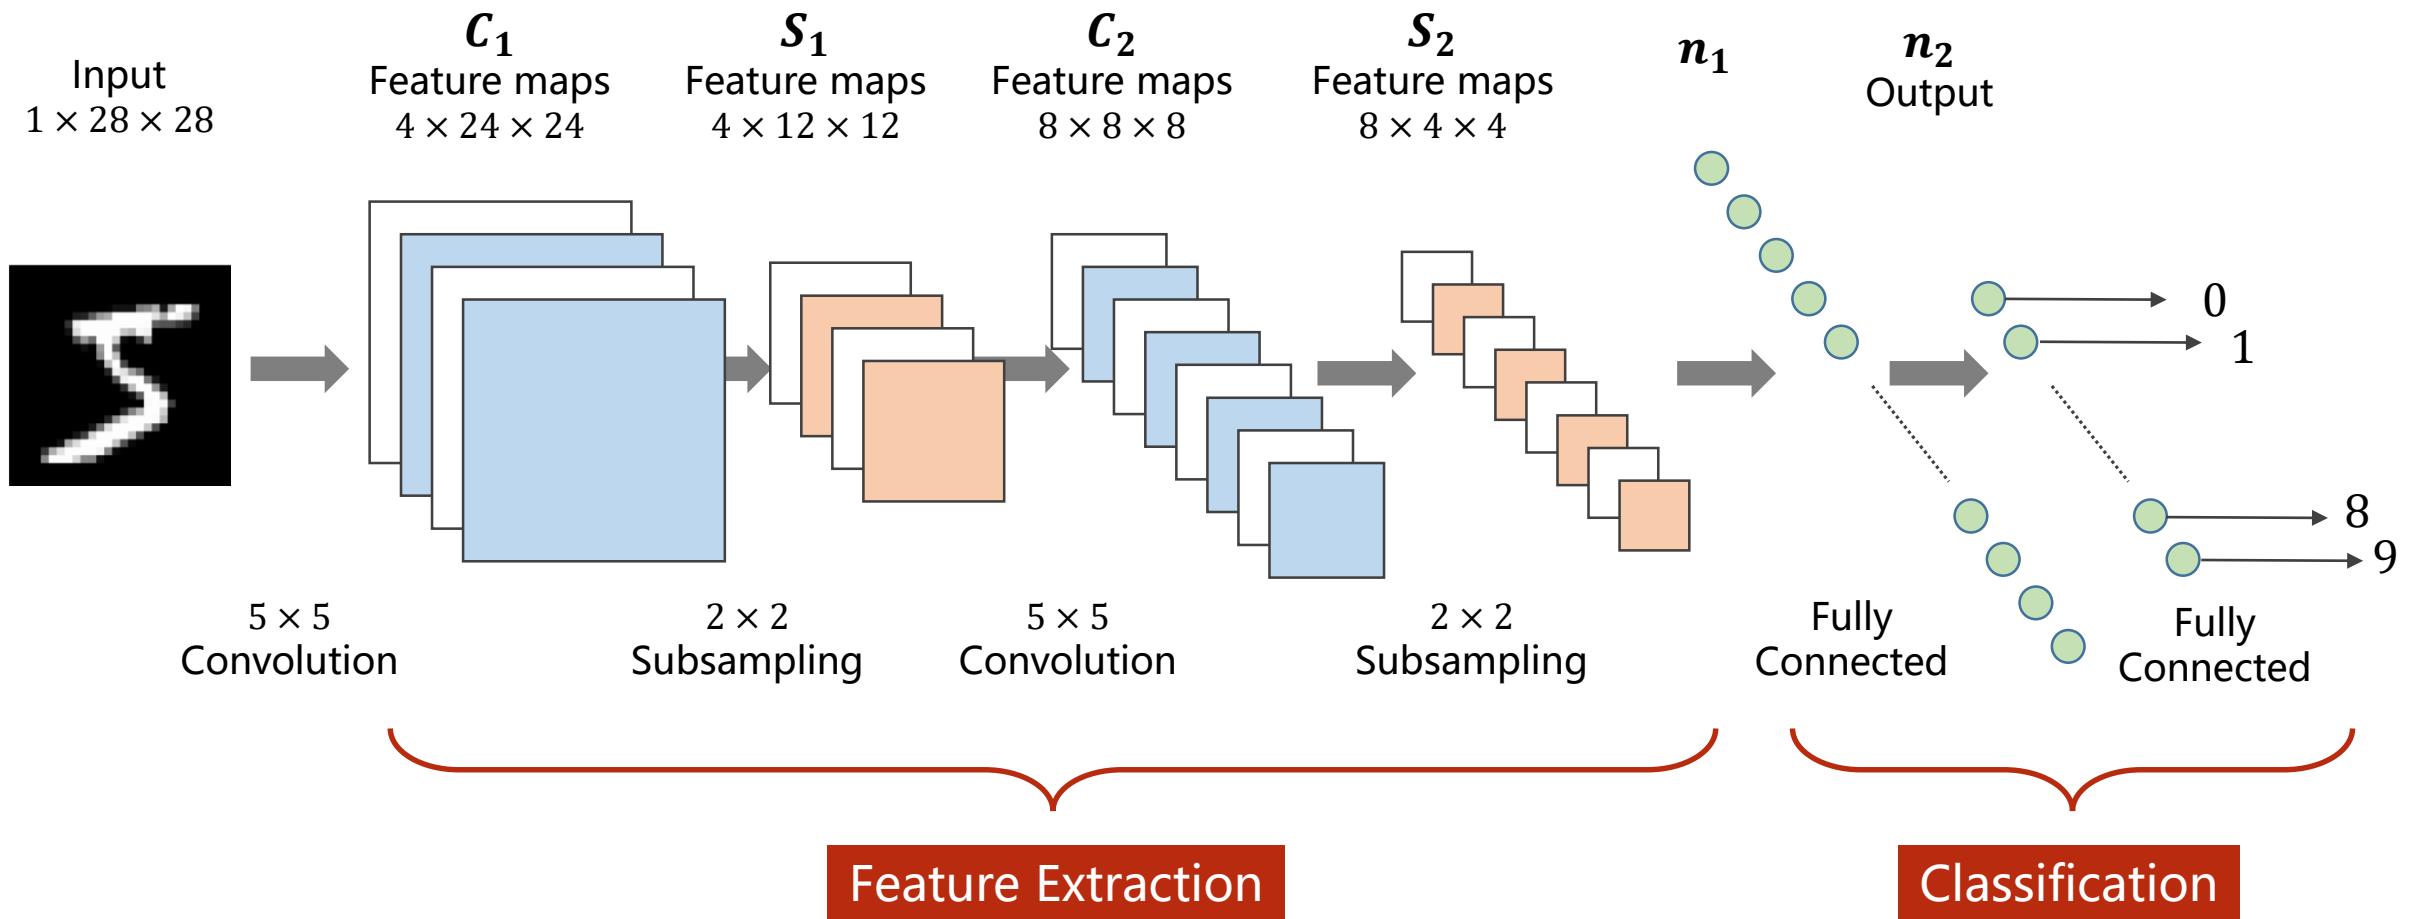
## 12. Basic RNN

循环神经网络：适合具有序列特征的输入数据

# Revision: DNN

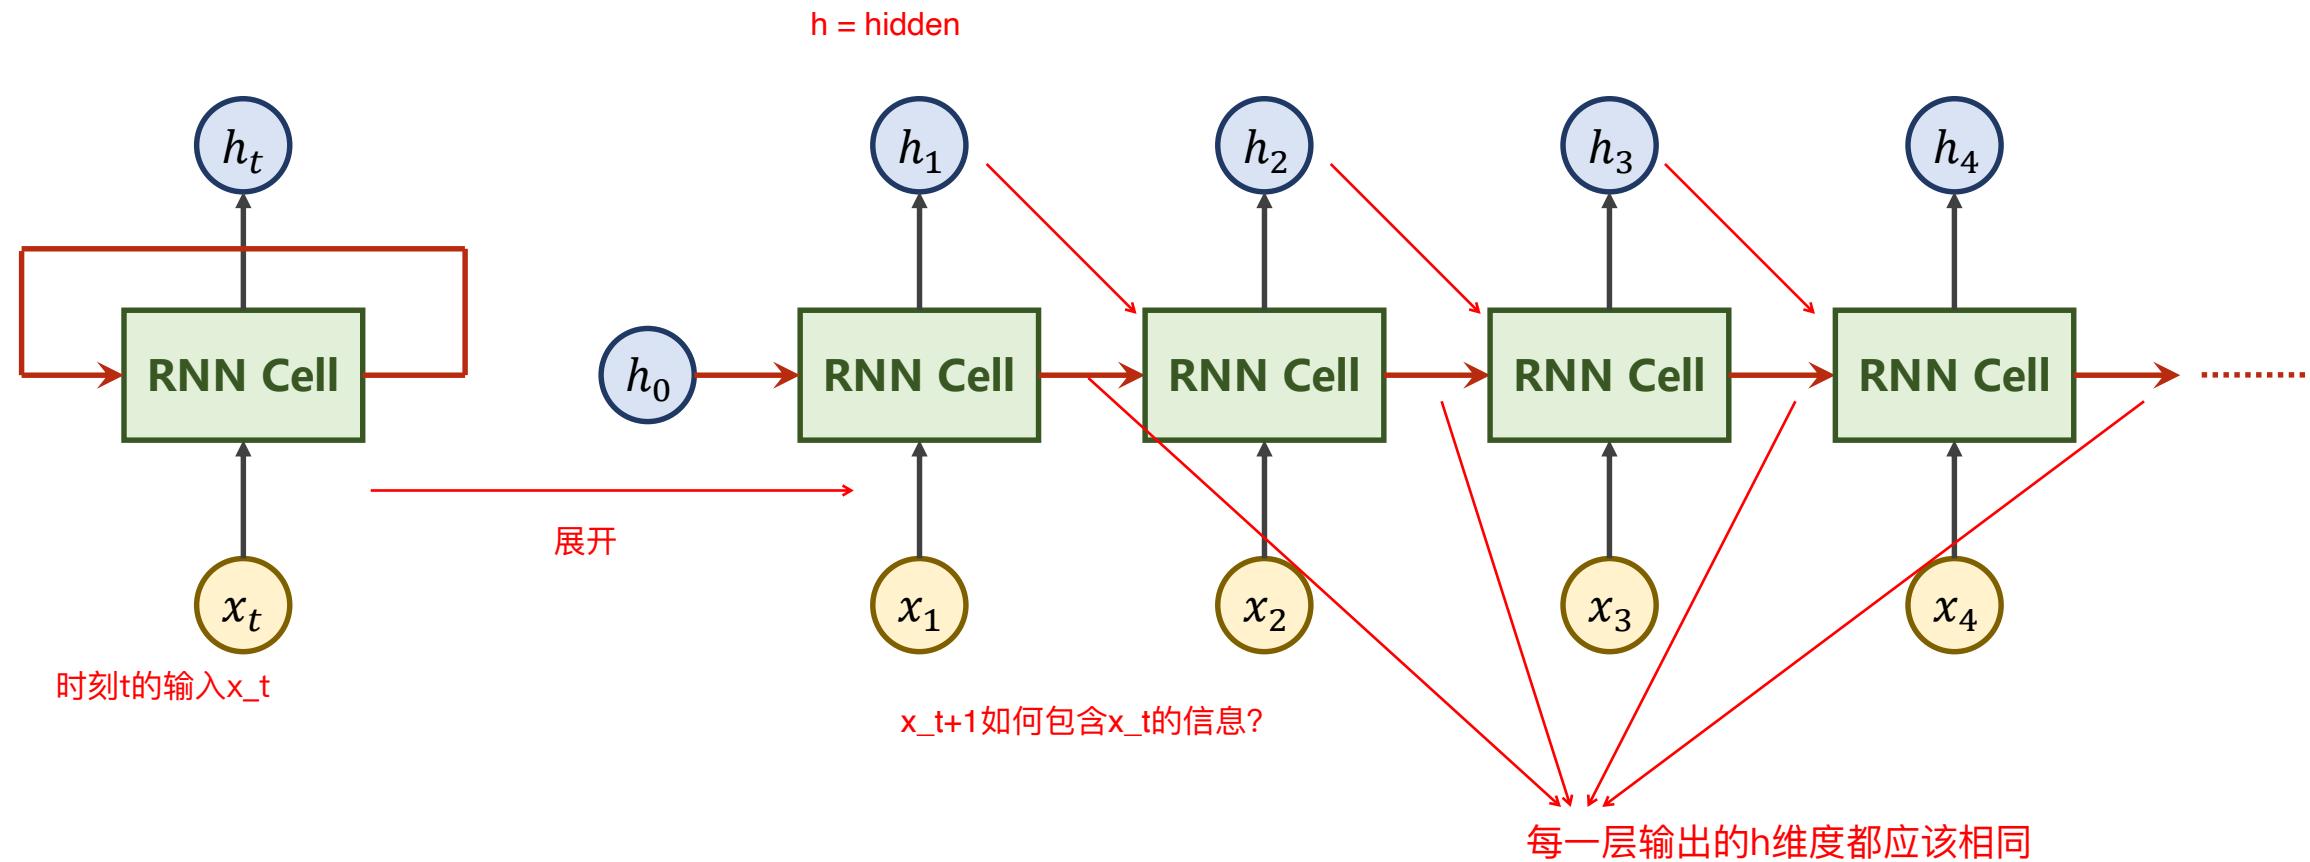


# Revision: CNN



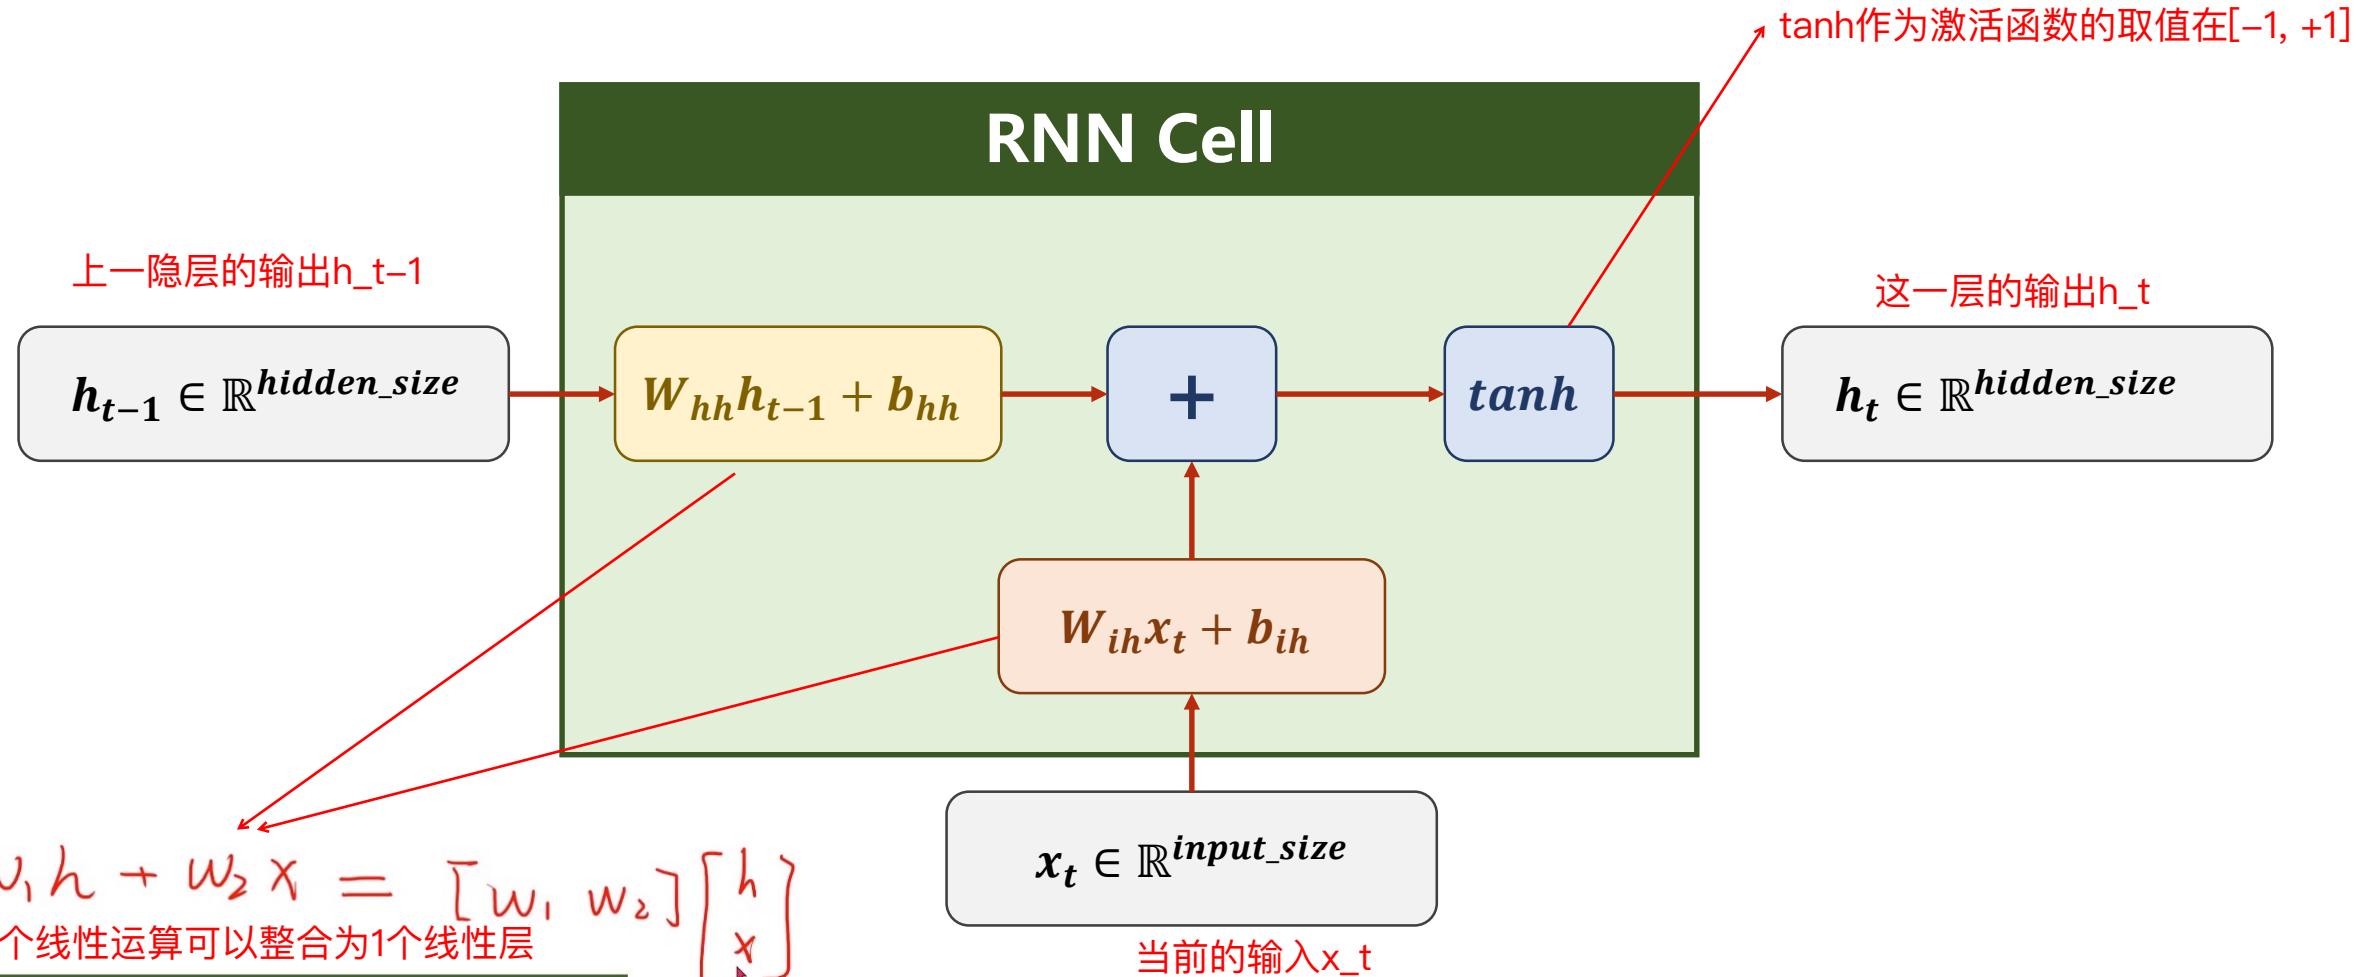
# What is RNNs?

这里的RNN Cell是循环共用的，其本质是一个线性层，但是其权重受序列整体的影响。

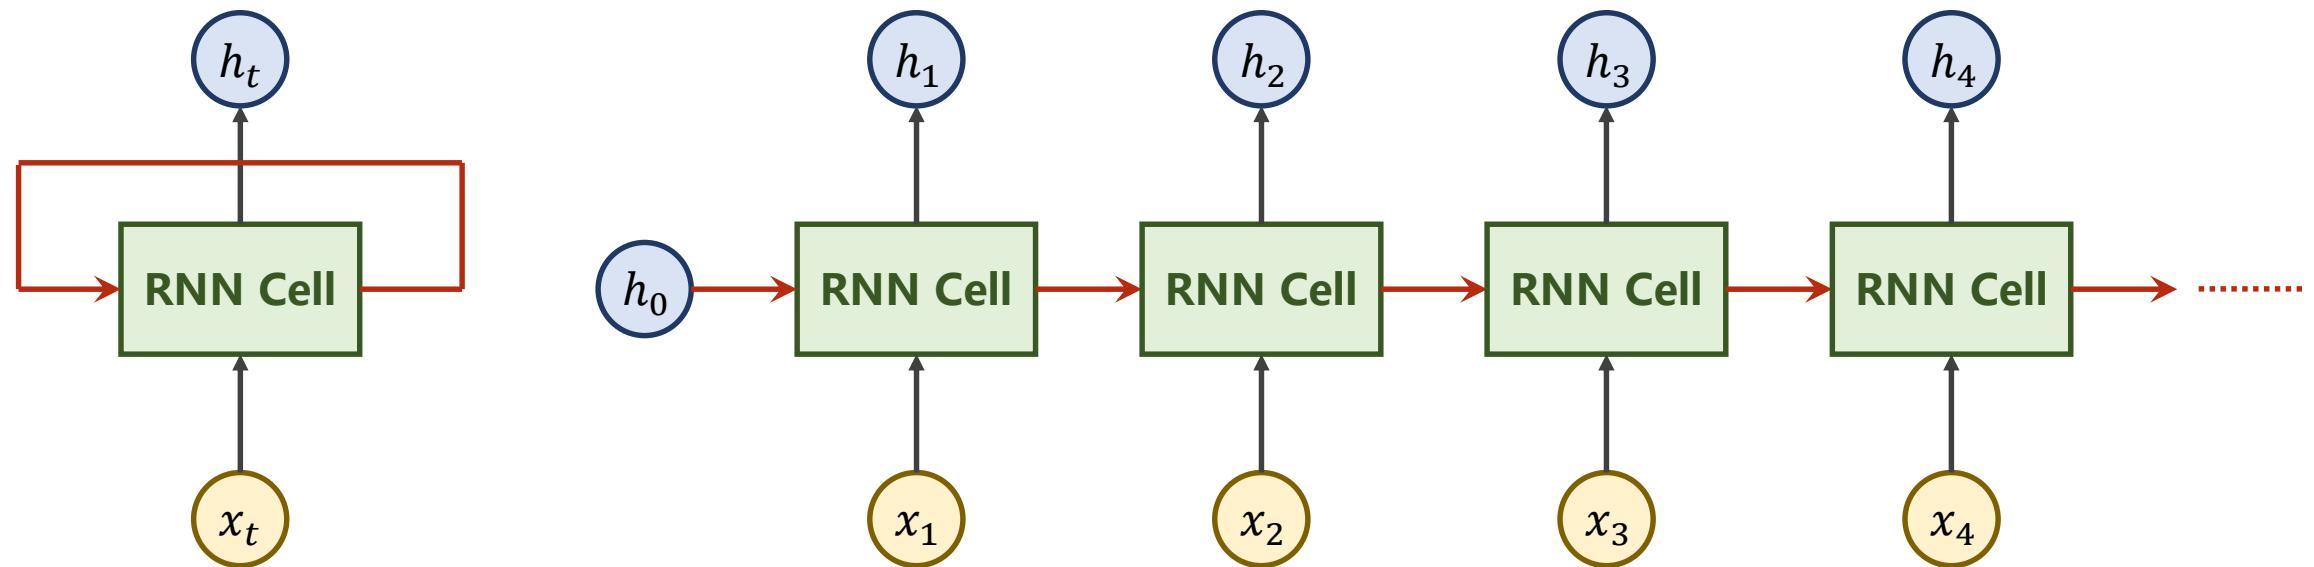


# What is RNN?

$$h_t = \tanh(W_{ih}x_t + b_{ih} + W_{hh}h_{t-1} + b_{hh})$$

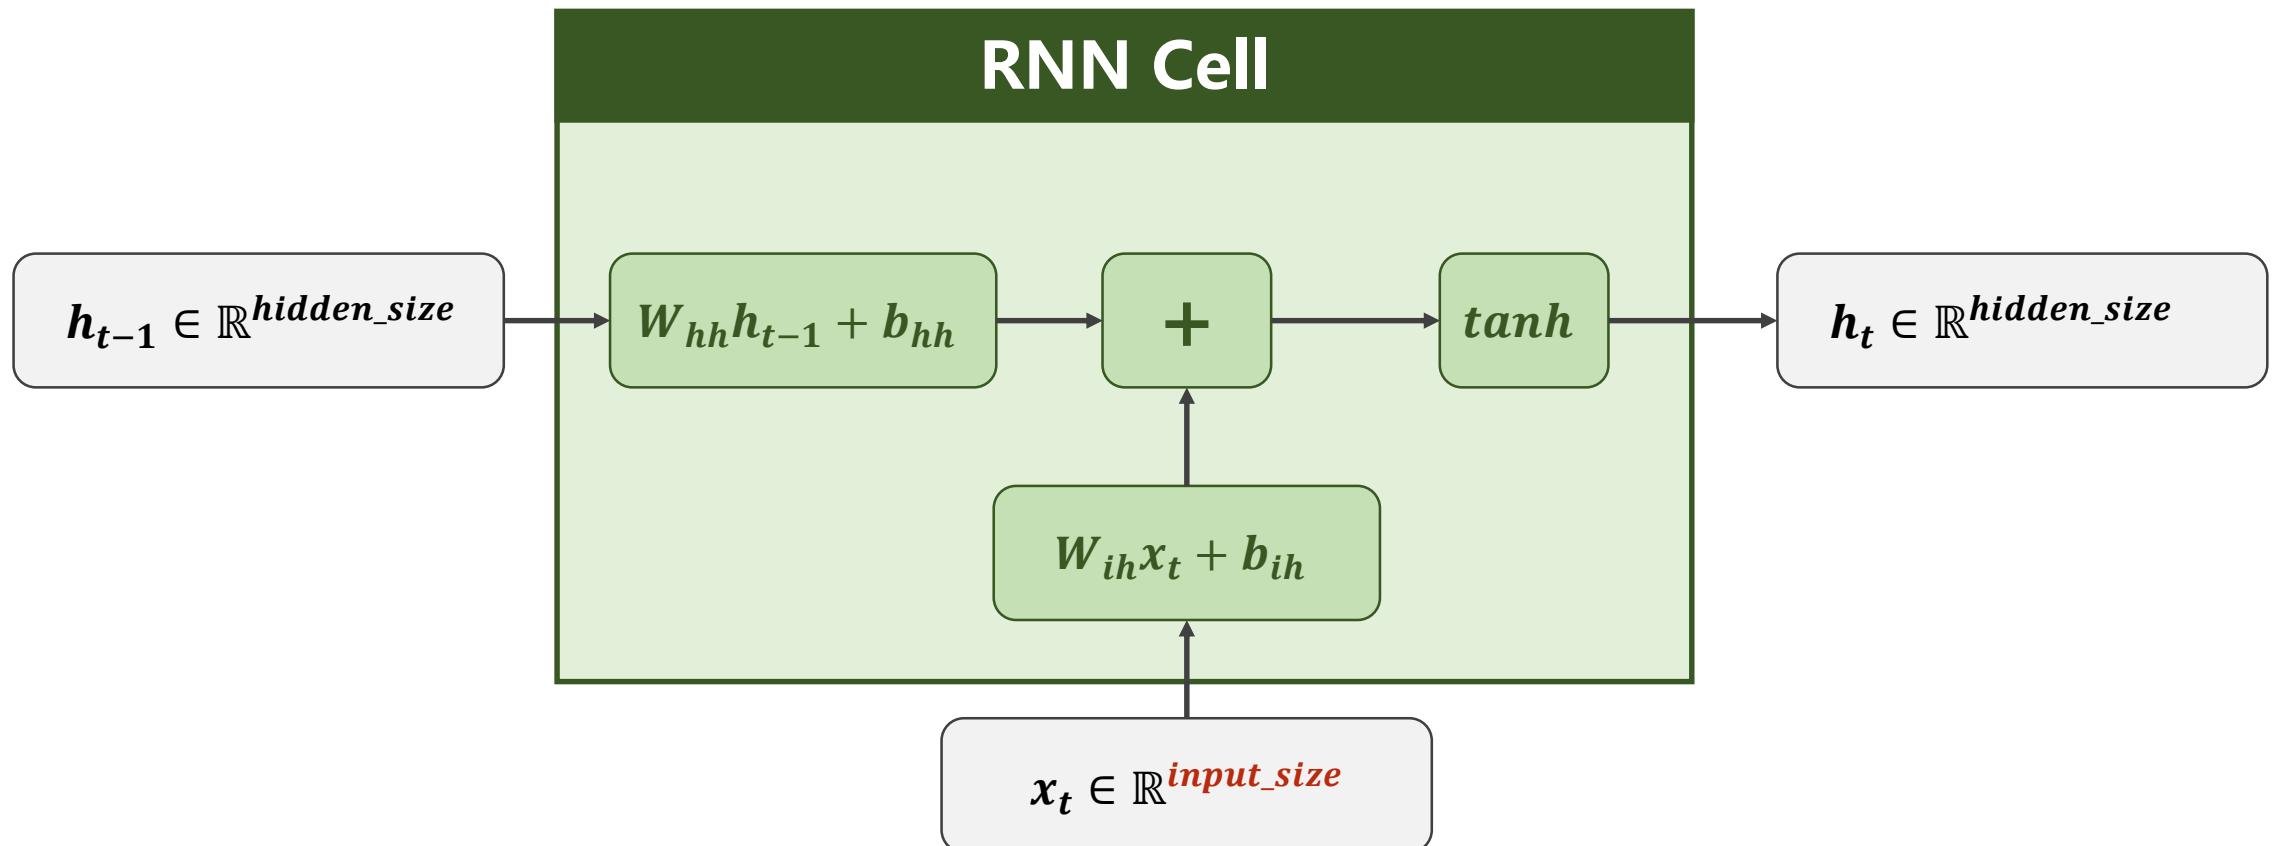


# What is RNNs?



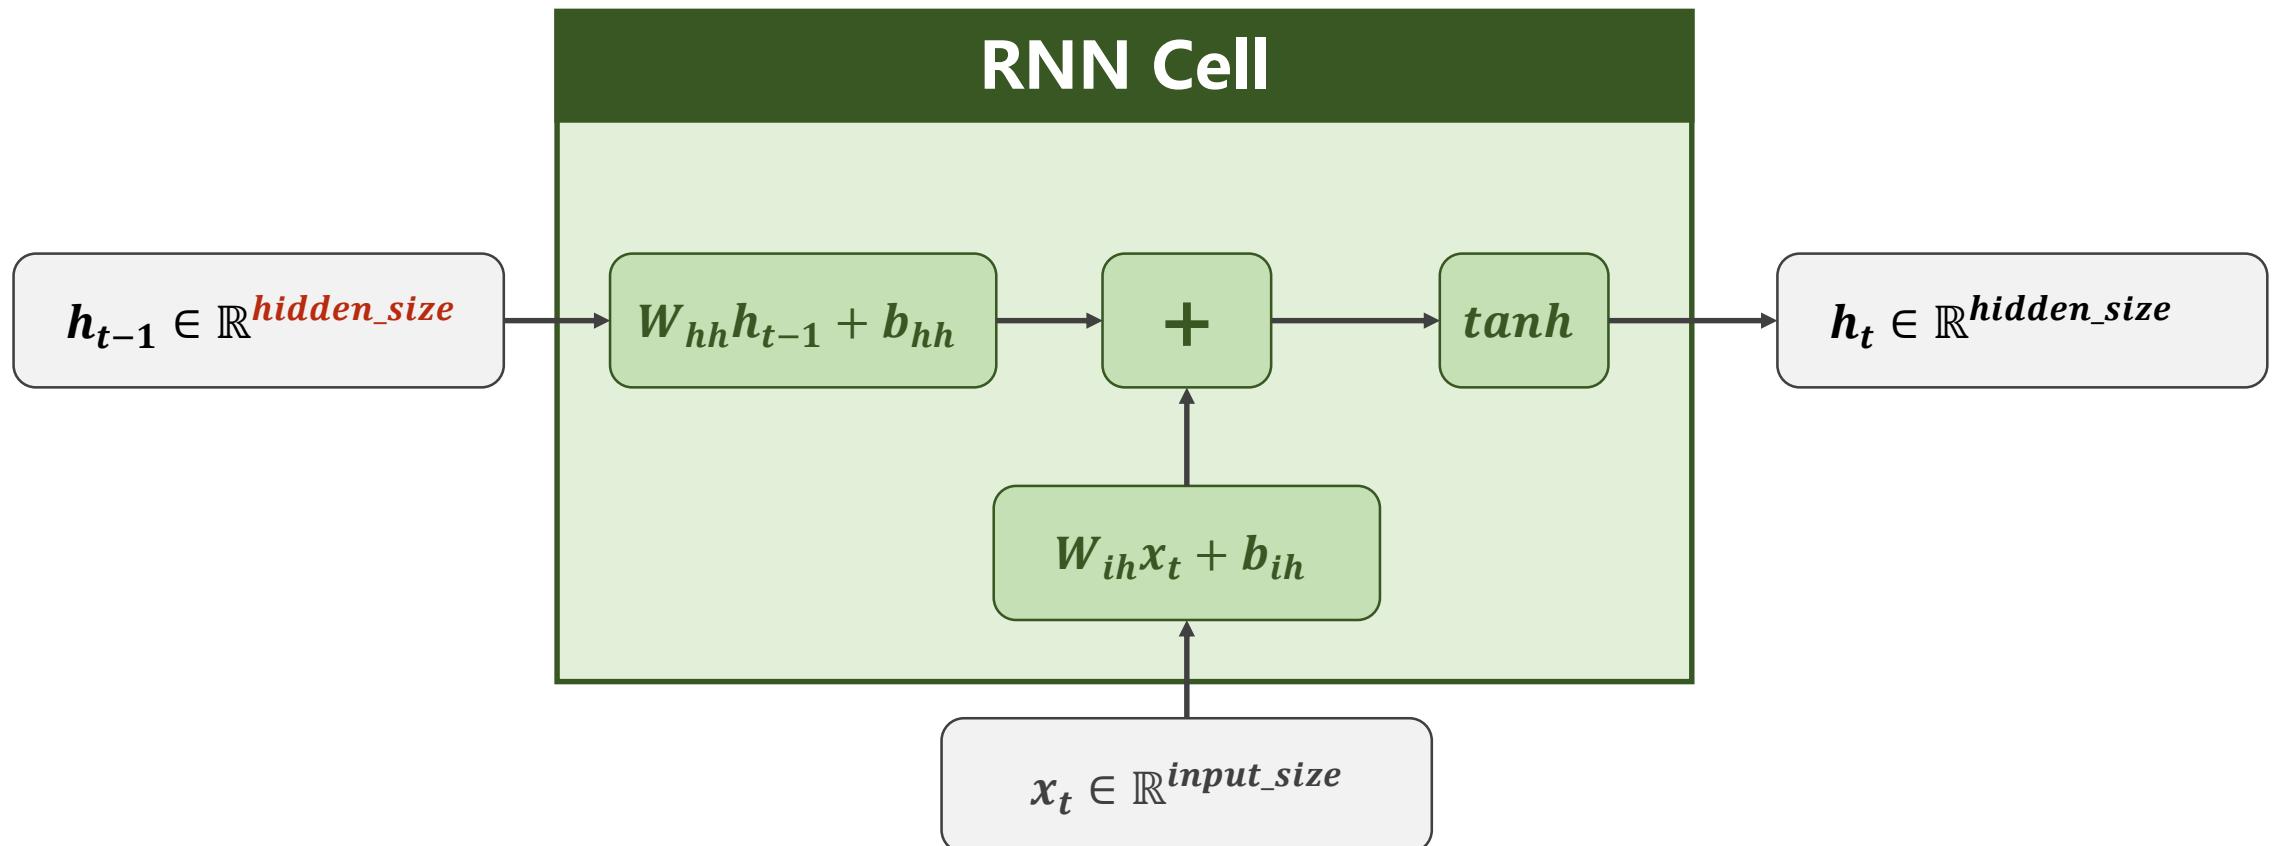
# RNN Cell in PyTorch

```
cell = torch.nn.RNNCell(input_size=input_size, hidden_size=hidden_size)
```



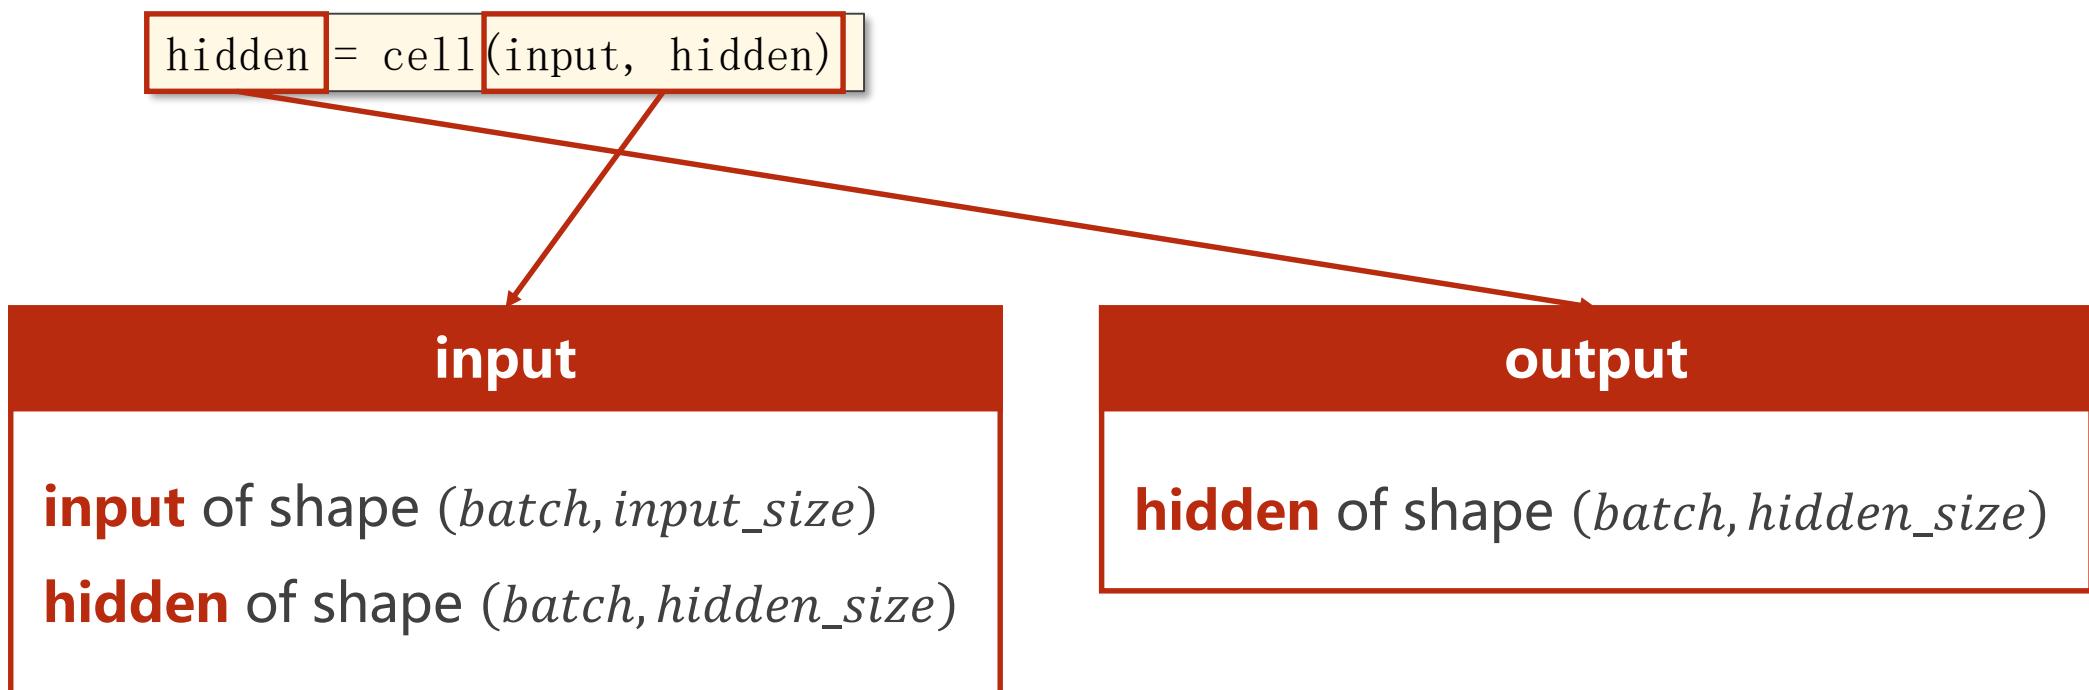
# RNN Cell in PyTorch

```
cell = torch.nn.RNNCell(input_size=input_size, hidden_size=hidden_size)
```



# RNN Cell in PyTorch

```
cell = torch.nn.RNNCell(input_size=input_size, hidden_size=hidden_size)
```



# How to use RNNCell

- Suppose we have sequence with below properties:
  - $batchSize = 1$
  - $seqLen = 3$
  - $inputSize = 4$
  - $hiddenSize = 2$

# How to use RNNCell

- Suppose we have sequence with below properties:
  - ***batchSize*** = 1
  - *seqLen* = 3
  - ***inputSize*** = 4
  - *hiddenSize* = 2
- So the shape of inputs and outputs of RNNCell:
  - ***input.shape*** = (***batchSize, inputSize***)
  - *output.shape* = (*batchSize, hiddenSize*)

# How to use RNNCell

- Suppose we have sequence with below properties:
  - ***batchSize*** = 1
  - *seqLen* = 3
  - *inputSize* = 4
  - ***hiddenSize*** = 2
- So the shape of inputs and outputs of RNNCell:
  - *input.shape* = (*batchSize*, *inputSize*)
  - ***output.shape*** = (***batchSize*, *hiddenSize***)

# How to use RNNCell

- Suppose we have sequence with below properties:
  - ***batchSize*** = 1
  - ***seqLen*** = 3
  - ***inputSize*** = 4
  - ***hiddenSize*** = 2
- So the shape of inputs and outputs of RNNCell:
  - *input.shape* = (*batchSize, inputSize*)
  - *output.shape* = (*batchSize, hiddenSize*)
- The sequence can be warped in one Tensor with shape:
  - *dataset.shape* = (***seqLen, batchSize, inputSize***)  
                          序列长度      N

# How to use RNNCell

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2 } Parameters

cell = torch.nn.RNNCell(input_size=input_size, hidden_size=hidden_size)

# (seq, batch, features)
dataset = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(batch_size, hidden_size)

for idx, input in enumerate(dataset):
    print('=' * 20, idx, '=' * 20)
    print('Input size: ', input.shape)

    hidden = cell(input, hidden)

    print('outputs size: ', hidden.shape)
    print(hidden)
```

# How to use RNNCell

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2

cell = torch.nn.RNNCell(input_size=input_size, hidden_size=hidden_size)

# (seq, batch, features)
dataset = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(batch_size, hidden_size)

for idx, input in enumerate(dataset):
    print('=' * 20, idx, '=' * 20)
    print('Input size: ', input.shape)

    hidden = cell(input, hidden)

    print('outputs size: ', hidden.shape)
    print(hidden)
```

Construction of RNNCell

# How to use RNNCell

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2

cell = torch.nn.RNNCell(input_size=input_size, hidden_size=hidden_size)

# (seq, batch, features)
dataset = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(batch_size, hidden_size)

for idx, input in enumerate(dataset):
    print('=' * 20, idx, '=' * 20)
    print('Input size: ', input.shape)

    hidden = cell(input, hidden)

    print('outputs size: ', hidden.shape)
    print(hidden)
```

Wrapping the sequence into:  
 $(seqLen, batchSize, inputSize)$

# How to use RNNCell

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2

cell = torch.nn.RNNCell(input_size=input_size, hidden_size=hidden_size)

# (seq, batch, features)
dataset = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(batch_size, hidden_size)

for idx, input in enumerate(dataset):
    print('=' * 20, idx, '=' * 20)
    print('Input size: ', input.shape)

    hidden = cell(input, hidden)

    print('outputs size: ', hidden.shape)
    print(hidden)
```

Initializing the **hidden** to zero

# How to use RNNCell

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2

cell = torch.nn.RNNCell(input_size=input_size, hidden_size=hidden_size)

# (seq, batch, features)
dataset = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(batch_size, hidden_size)

for idx, input in enumerate(dataset):
    print('=' * 20, idx, '=' * 20)
    print('Input size: ', input.shape)

    hidden = cell(input, hidden)

    print('outputs size: ', hidden.shape)
    print(hidden)
```

The shape of input is:  
*(batchSize, inputSize)*

Input size: torch.Size([1, 4])

# How to use RNNCell

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2

cell = torch.nn.RNNCell(input_size=input_size, hidden_size=hidden_size)

# (seq, batch, features)
dataset = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(batch_size, hidden_size)

for idx, input in enumerate(dataset):
    print('=' * 20, idx, '=' * 20)
    print('Input size: ', input.shape)

    hidden = cell(input, hidden)

    print('hidden size: ', hidden.shape)
    print(hidden)
```

The shape of **hidden** is:  
**(batchSize, hiddenSize)**

hidden size: torch.Size([1, 2])

# How to use RNNCell

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2

cell = torch.nn.RNNCell(input_size=input_size, hidden_size=hidden_size)

# (seq, batch, features)
dataset = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(batch_size, hidden_size)

for idx, input in enumerate(dataset):
    print('=' * 20, idx, '=' * 20)
    print('Input size: ', input.shape)

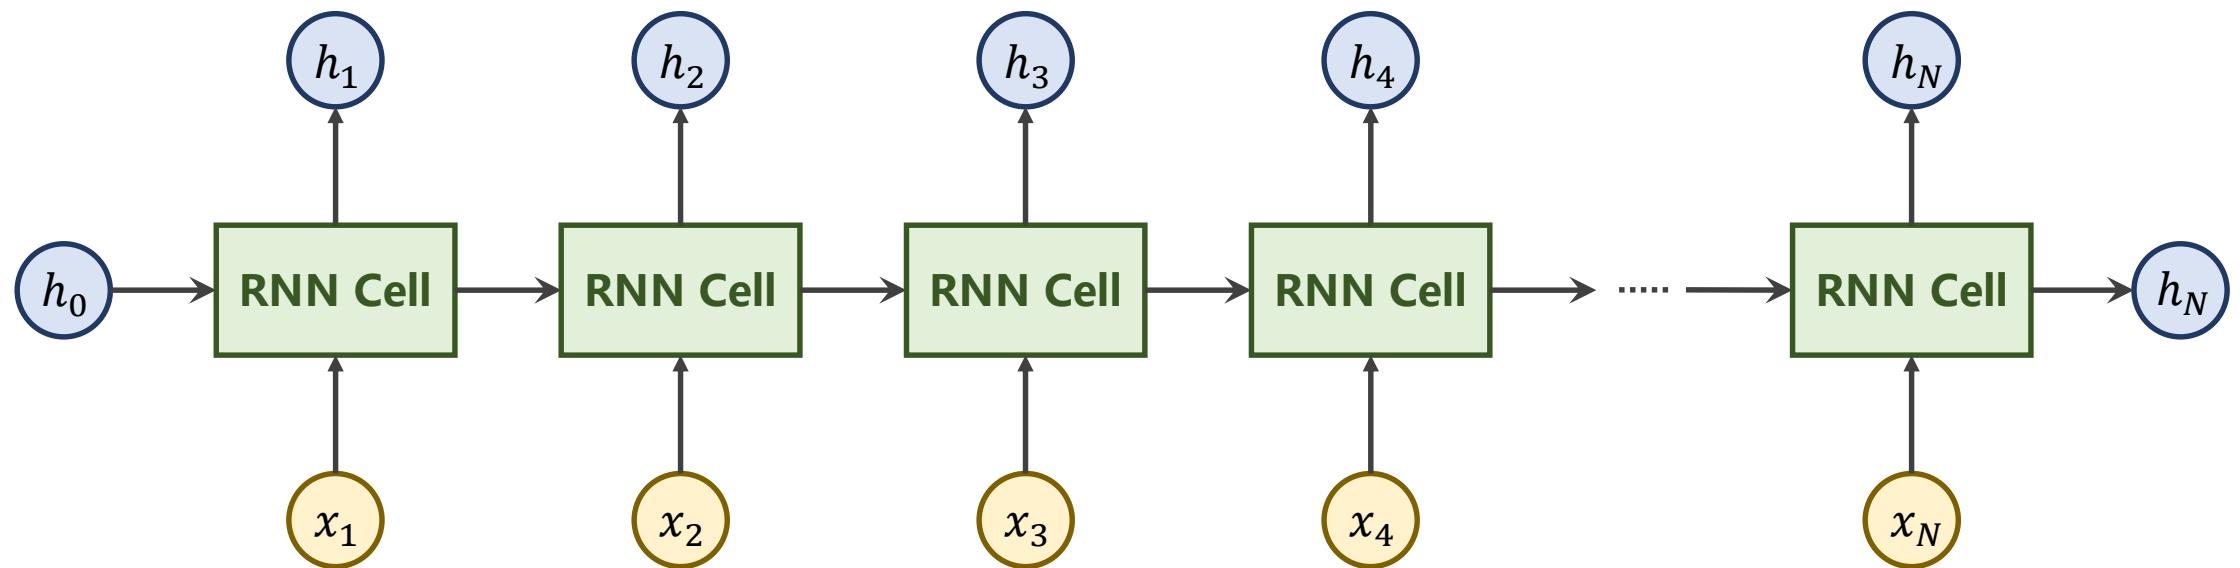
    hidden = cell(input, hidden)

    print('outputs size: ', hidden.shape)
    print(hidden)
```

```
===== 0 =====
Input size: torch.Size([1, 4])
outputs size: torch.Size([1, 2])
tensor([-0.1579,  0.5140])
===== 1 =====
Input size: torch.Size([1, 4])
outputs size: torch.Size([1, 2])
tensor([-0.9577,  0.6502])
===== 2 =====
Input size: torch.Size([1, 4])
outputs size: torch.Size([1, 2])
tensor([-0.7661, -0.9128])
```

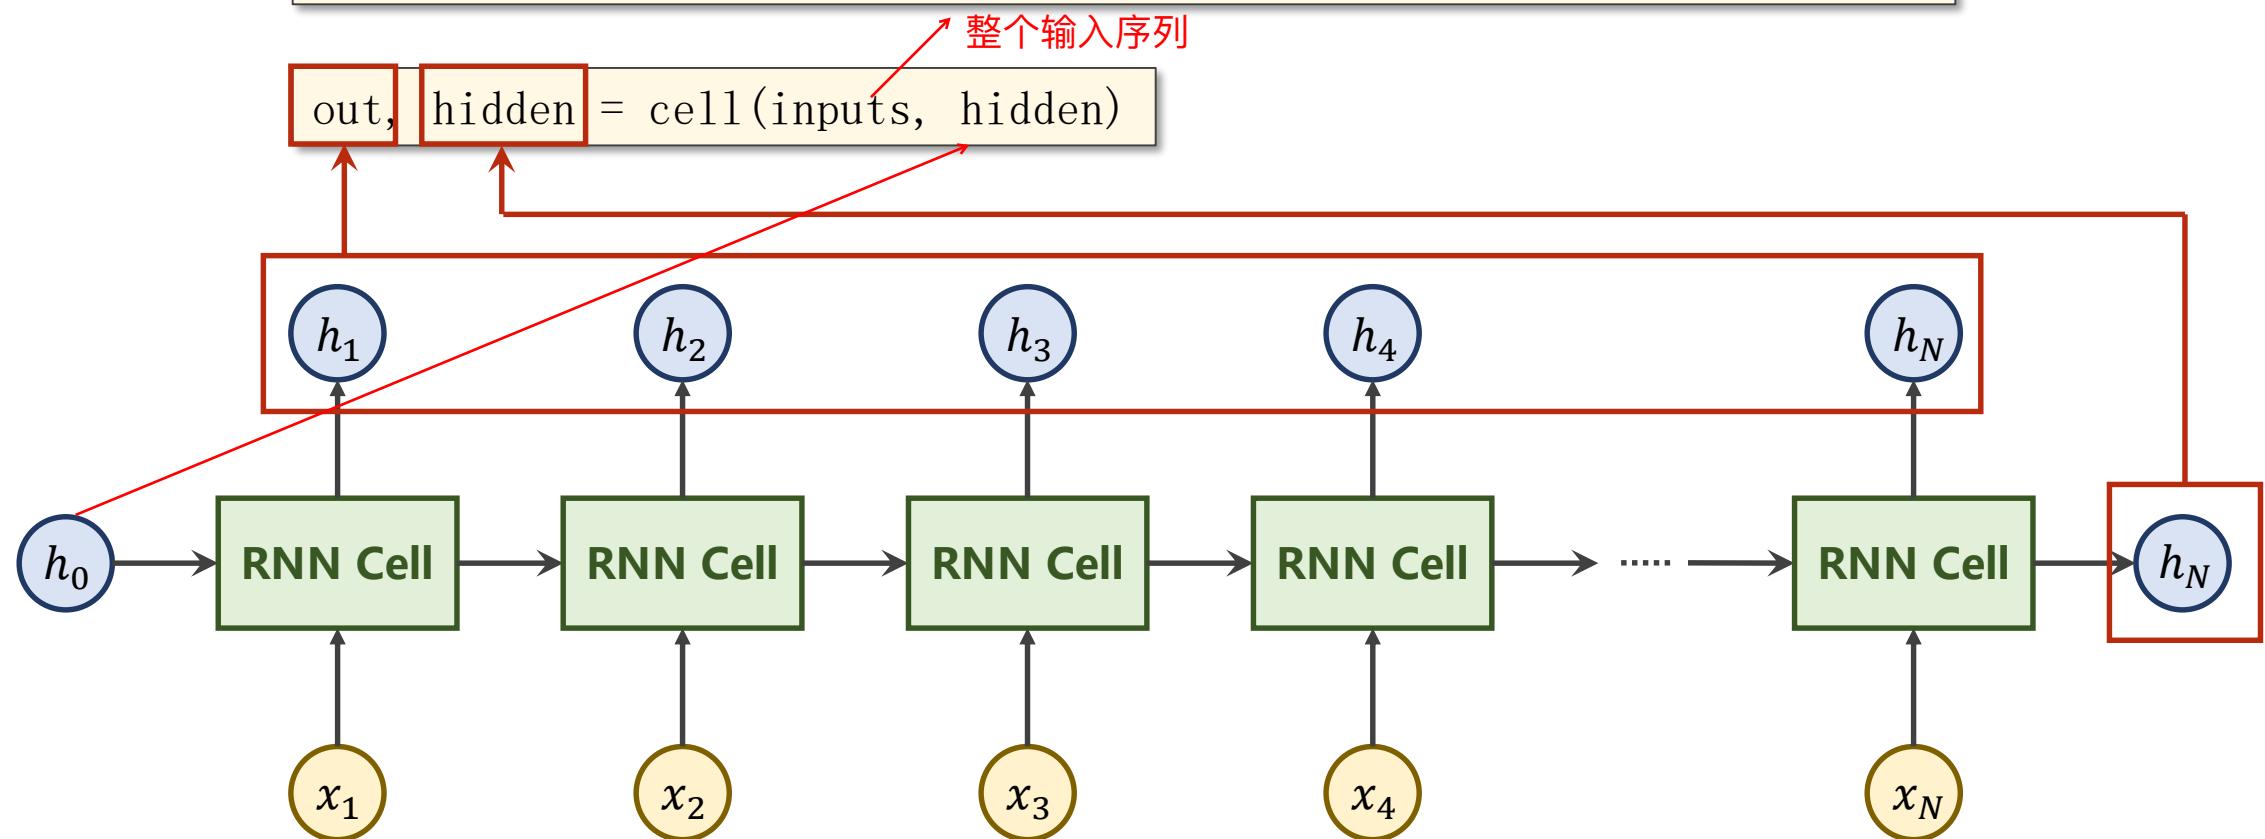
# How to use RNN

```
cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,  
                     num_layers=num_layers)
```



# How to use RNN

```
cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,  
                     num_layers=num_layers) RNN的层数
```



# How to use RNN

```
cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,  
                    num_layers=num_layers)
```

```
out, hidden = cell(inputs, hidden)
```

**input**

**input** of shape (*seqSize, batch, input\_size*)  
**hidden** of shape (*numLayers, batch, hidden\_size*)

**output**

**Output** of shape (*seqLen, batch, hidden\_size*)  
**hidden** of shape (*numLayers, batch, hidden\_size*)

# How to use RNN

- Suppose we have sequence with below properties:
  - *batchSize*
  - *seqLen*
  - *inputSize, hiddenSize,*
  - *numLayers*

# How to use RNN

- Suppose we have sequence with below properties:
  - ***batchSize***
  - ***seqLen***
  - ***inputSize, hiddenSize,***
  - ***numLayers***
- The shape of ***input*** and ***h\_0*** of RNN:
  - $input.shape = (\text{seqLen}, \text{batchSize}, \text{inputSize})$
  - $h_0.shape = (\text{numLayers}, \text{batchSize}, \text{hiddenSize})$

# How to use RNN

- Suppose we have sequence with below properties:
  - ***batchSize***
  - *seqLen*
  - *inputSize*, ***hiddenSize***,
  - ***numLayers***
- The shape of ***input*** and ***h\_0*** of RNN:
  - *input.shape* = (*seqLen*, *batchSize*, *inputSize*)
  - *h\_0.shape* = (***numLayers***, ***batchSize***, ***hiddenSize***)

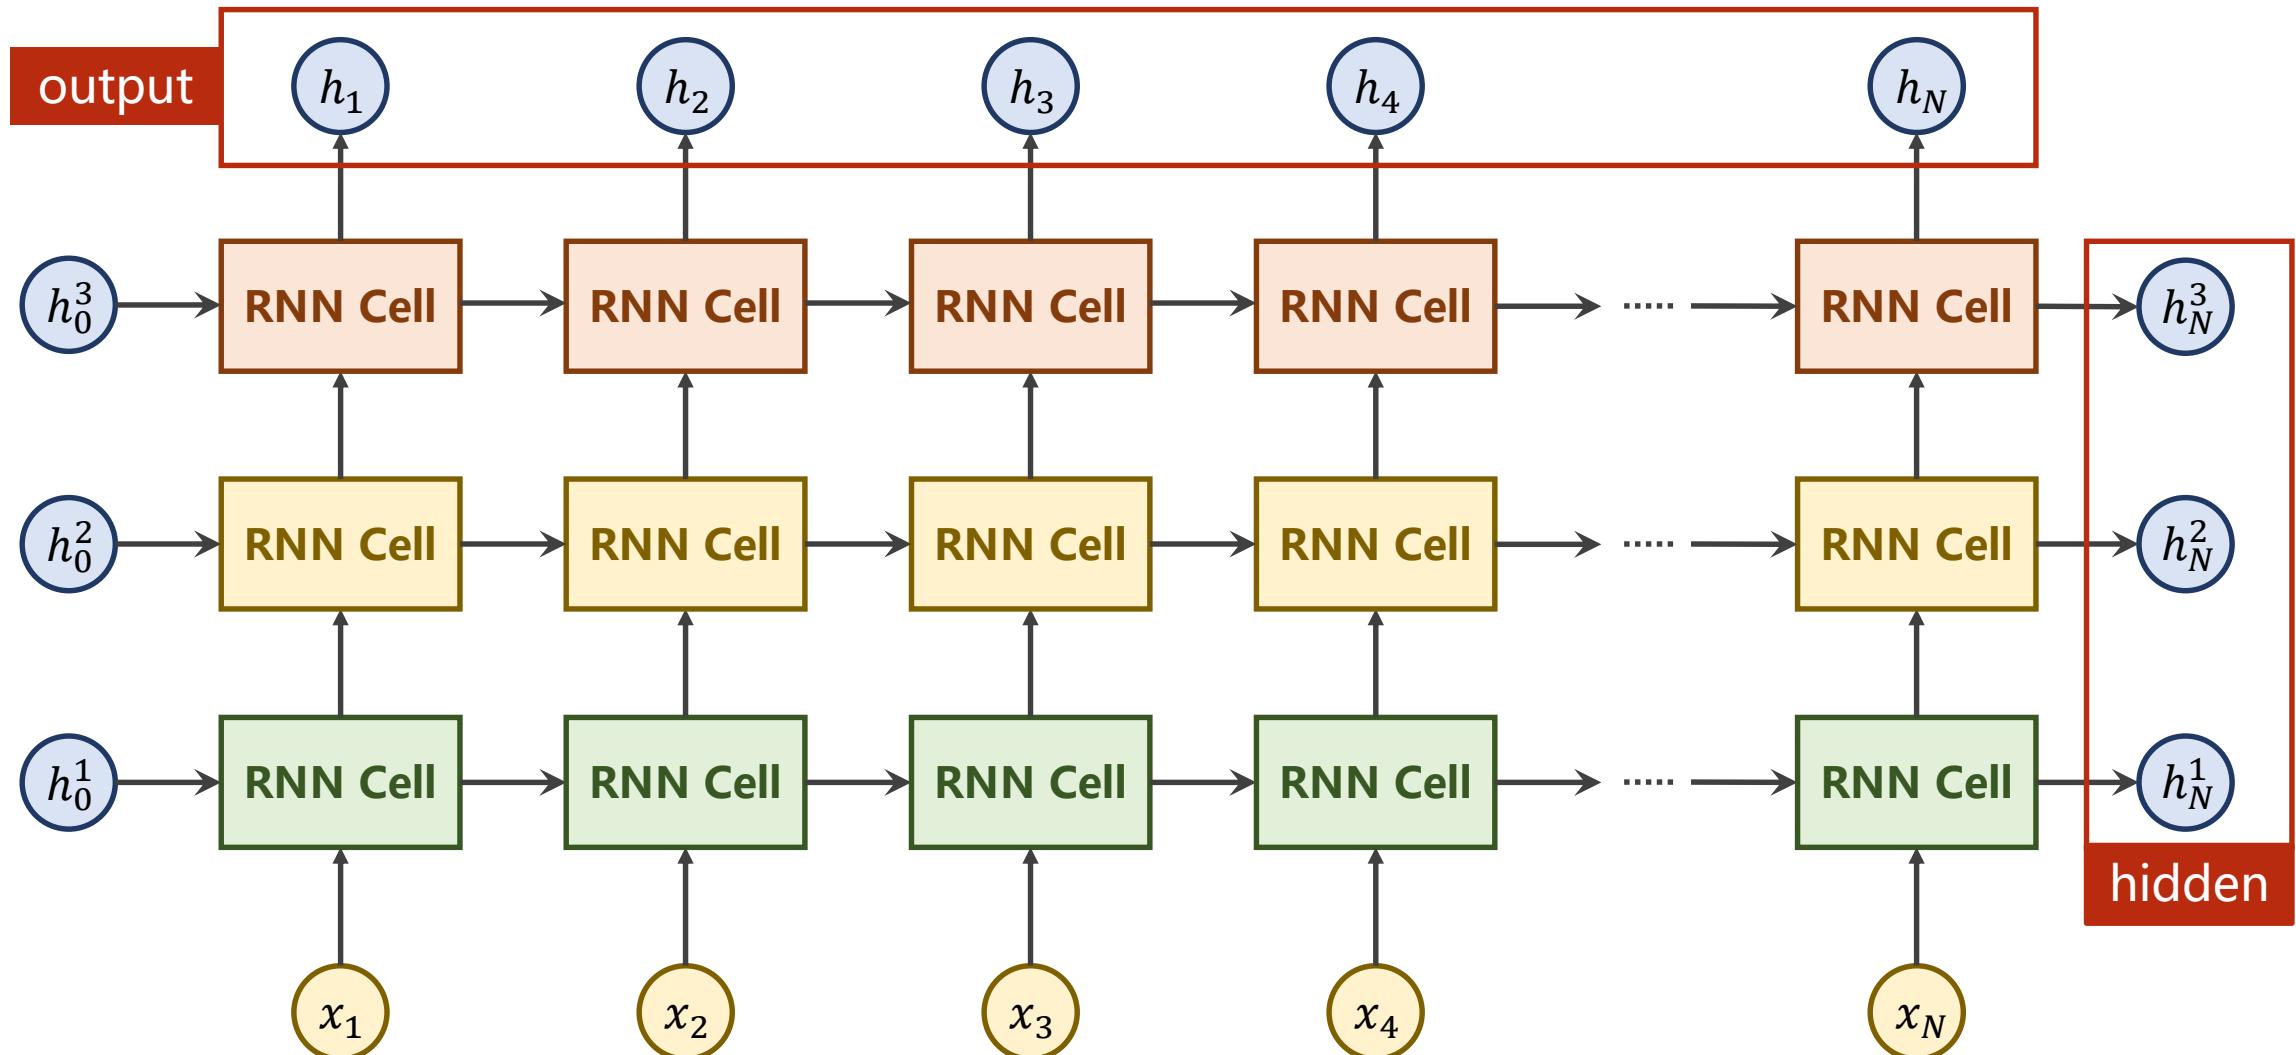
# How to use RNN

- Suppose we have sequence with below properties:
  - ***batchSize***
  - ***seqLen***
  - ***inputSize, hiddenSize***,
  - ***numLayers***
- The shape of ***input*** and ***h\_0*** of RNN:
  - $\text{input.shape} = (\text{seqLen}, \text{batchSize}, \text{inputSize})$
  - $\text{h\_0.shape} = (\text{numLayers}, \text{batchSize}, \text{hiddenSize})$
- The shape of ***output*** and ***h\_n*** of RNN:
  - $\text{output.shape} = (\text{seqLen}, \text{batchSize}, \text{hiddenSize})$
  - $\text{h\_n.shape} = (\text{numLayers}, \text{batchSize}, \text{hiddenSize})$

# How to use RNN

- Suppose we have sequence with below properties:
  - **batchSize**
  - *seqLen*
  - *inputSize*, **hiddenSize**,
  - **numLayers**
- The shape of **input** and ***h\_0*** of RNN:
  - *input.shape* = (*seqLen*, *batchSize*, *inputSize*)
  - *h\_0.shape* = (*numLayers*, *batchSize*, *hiddenSize*)
- The shape of **output** and ***h\_n*** of RNN:
  - *output.shape* = (*seqLen*, *batchSize*, *hiddenSize*)
  - *h\_n.shape* = (**numLayers**, **batchSize**, **hiddenSize**)

# How to use RNN - numLayers



# How to use RNN

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2
num_layers = 1 } Parameters

cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,
                     num_layers=num_layers)

# (seqLen, batchSize, inputSize)
inputs = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(num_layers, batch_size, hidden_size)

out, hidden = cell(inputs, hidden)

print('Output size:', out.shape)
print('Output:', out)
print('Hidden size:', hidden.shape)
print('Hidden:', hidden)
```

# How to use RNN

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2
num_layers = 1

cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,
                    num_layers=num_layers)

# (seqLen, batchSize, inputSize)
inputs = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(num_layers, batch_size, hidden_size)

out, hidden = cell(inputs, hidden)

print('Output size:', out.shape)
print('Output:', out)
print('Hidden size:', hidden.shape)
print('Hidden:', hidden)
```

## Construction of RNN

# How to use RNN

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2
num_layers = 1

cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,
                     num_layers=num_layers)

# (seqLen, batchSize, inputSize)
inputs = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(num_layers, batch_size, hidden_size)

out, hidden = cell(inputs, hidden)

print('Output size:', out.shape)
print('Output:', out)
print('Hidden size:', hidden.shape)
print('Hidden:', hidden)
```

Wrapping the sequence into:  
 $(seqLen, batchSize, inputSize)$

# How to use RNN

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2
num_layers = 1

cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,
                    num_layers=num_layers)

# (seqLen, batchSize, inputSize)
inputs = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(num_layers, batch_size, hidden_size)

out, hidden = cell(inputs, hidden)

print('Output size:', out.shape)
print('Output:', out)
print('Hidden size:', hidden.shape)
print('Hidden:', hidden)
```

Initializing the **hidden** to zero

# How to use RNN

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2
num_layers = 1

cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,
                     num_layers=num_layers)

# (seqLen, batchSize, inputSize)
inputs = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(num_layers, batch_size, hidden_size)

out, hidden = cell(inputs, hidden)

print('Output size:', out.shape)
print('Output:', out)
print('Hidden size:', hidden.shape)
print('Hidden:', hidden)
```

The shape of *output* is:  
*(seqSize, batchSize, hiddenSize)*

# How to use RNN

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2
num_layers = 1

cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,
                     num_layers=num_layers)

# (seqLen, batchSize, inputSize)
inputs = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(num_layers, batch_size, hidden_size)

out, hidden = cell(inputs, hidden)

print('Output size:', out.shape)
print('Output:', out)
print('Hidden size:', hidden.shape)
print('Hidden:', hidden)
```

The shape of **hidden** is:  
**(numLayers, batchSize, hiddenSize)**

# How to use RNN

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2
num_layers = 1

cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,
                    num_layers=num_layers)

# (seqLen, batchSize, inputSize)
inputs = torch.randn(seq_len, batch_size, input_size)
hidden = torch.zeros(num_layers, batch_size, hidden_size)

out, hidden = cell(inputs, hidden)

print('Output size:', out.shape)
print('Output:', out)
print('Hidden size:', hidden.shape)
print('Hidden:', hidden)
```

Output size: torch.Size([3, 1, 2])  
Output: tensor([  
 [[ 0.2360, -0.5550]],  
 [[ 0.2685, 0.0106]],  
 [[ 0.1172, -0.2262]]])  
Hidden size: torch.Size([1, 1, 2])  
Hidden: tensor([[[ 0.1172, -0.2262]]])

# How to use RNN

- **batch\_first**: if True, the input and output tensors are provided as:
  - $(batchSize, seqLen, input\_size)$

```
cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,  
                     num_layers=num_layers, batch_first=True)
```

```
inputs = torch.randn(seq_len, batch_size, input_size)
```



```
inputs = torch.randn(batch_size, seq_len, input_size)
```

# How to use RNN

```
import torch

batch_size = 1
seq_len = 3
input_size = 4
hidden_size = 2
num_layers = 1

cell = torch.nn.RNN(input_size=input_size, hidden_size=hidden_size,
                     num_layers=num_layers, batch_first=True)

# (seqLen, batchSize, inputSize)
inputs = torch.randn(batch_size, seq_len, input_size)
hidden = torch.zeros(num_layers, batch_size, hidden_size)

out, hidden = cell(inputs, hidden)

print('Output size:', out.shape)
print('Output:', out)
print('Hidden size:', hidden.shape)
print('Hidden:', hidden)
```

Output size: torch.Size([1, 3, 2])

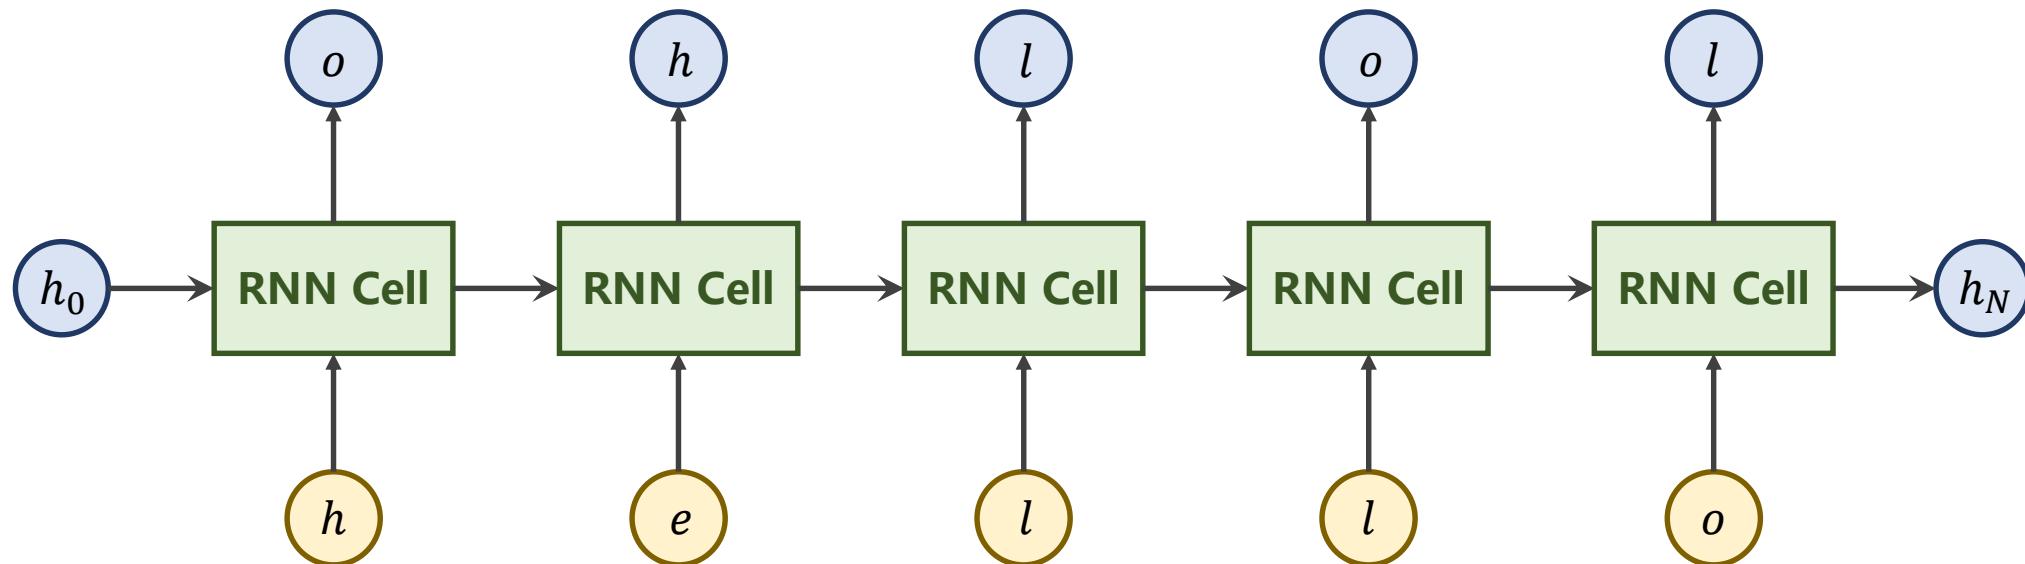
Output: tensor([[  
 [ 0.6582, -0.7281],  
 [ 0.5021, -0.8324],  
 [-0.5225, -0.5621]]])

Hidden size: torch.Size([1, 1, 2])

Hidden: tensor([[-0.5225, -0.5621]]])

# Example 12-1: Using RNNCell

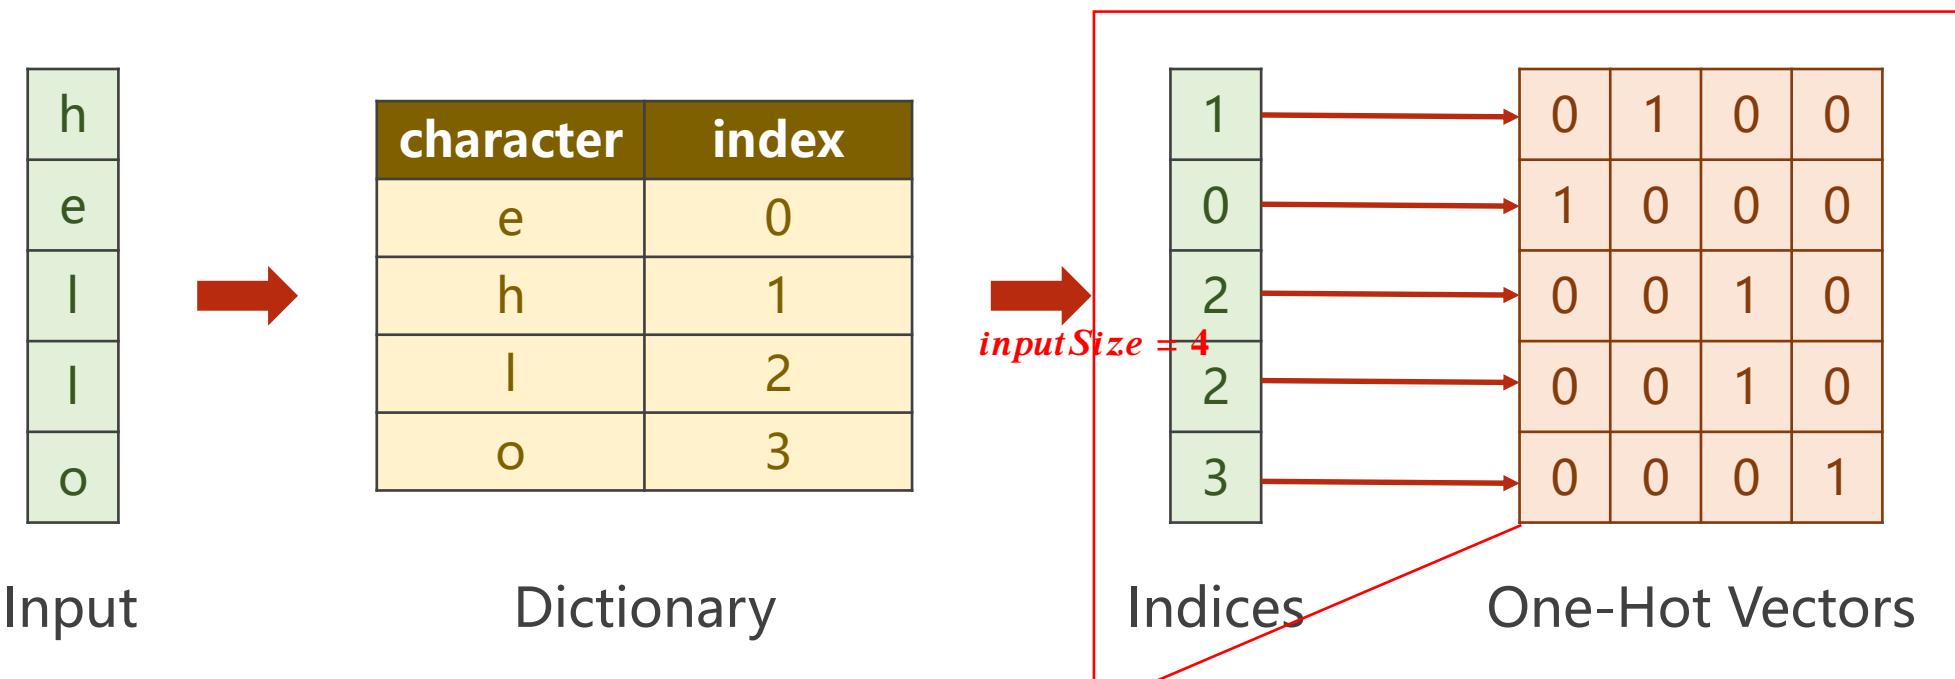
- Train a model to learn:
  - “hello” -> “ohlol” 训练RNN网络来实现序列的转换



# Example 12-1: Using RNNCell

- The inputs of RNN Cell should be vectors of numbers.

RNN的输入不能是字符串，应该是向量。首先构造词典，然后将字符串使用one-hot编码为向量。

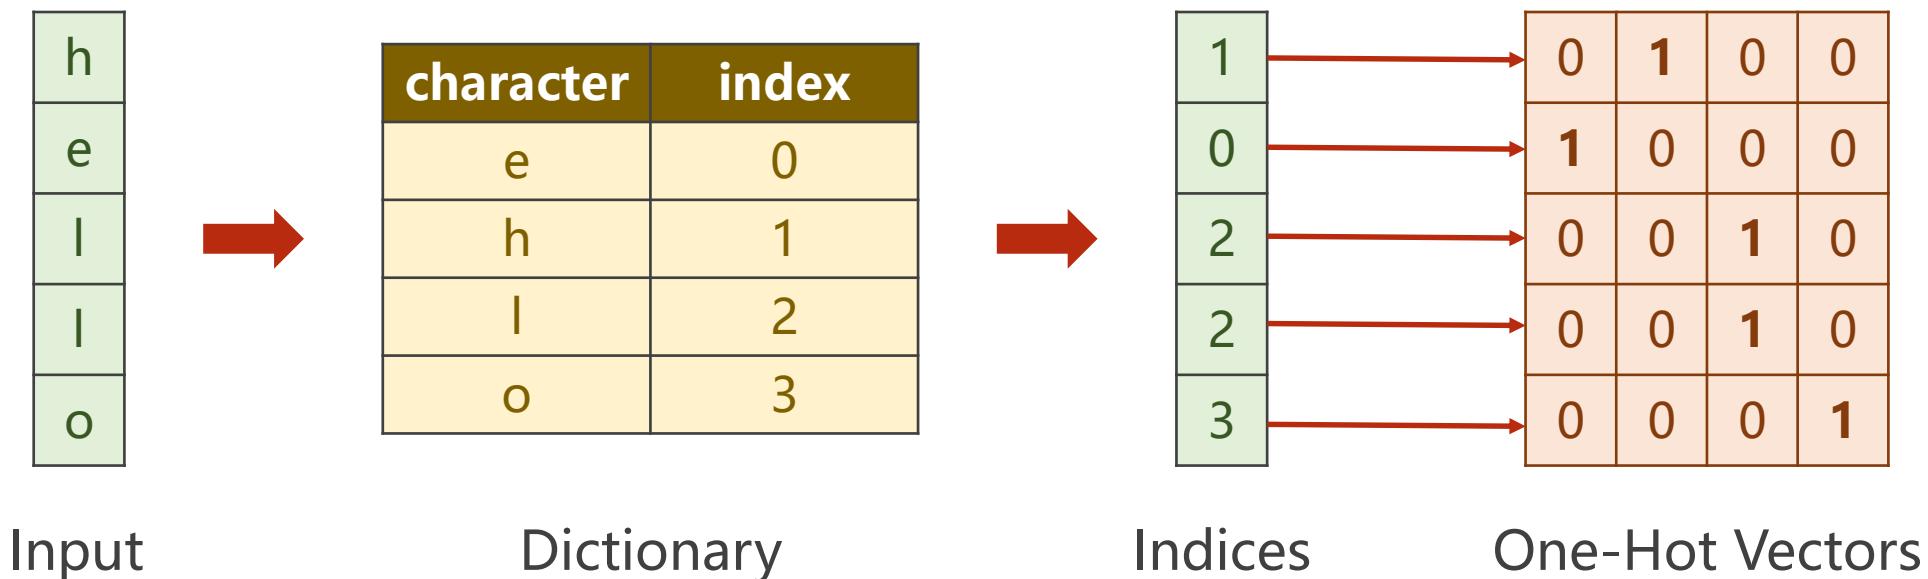


长度为n的包含m个字符的字符串，  
经过独热编码之后的向量大小为n\*m

One-Hot Vectors的宽度与词典的长度一致  
One-Hot Vectors中只有0和1

# Example 12-1: Using RNNCell

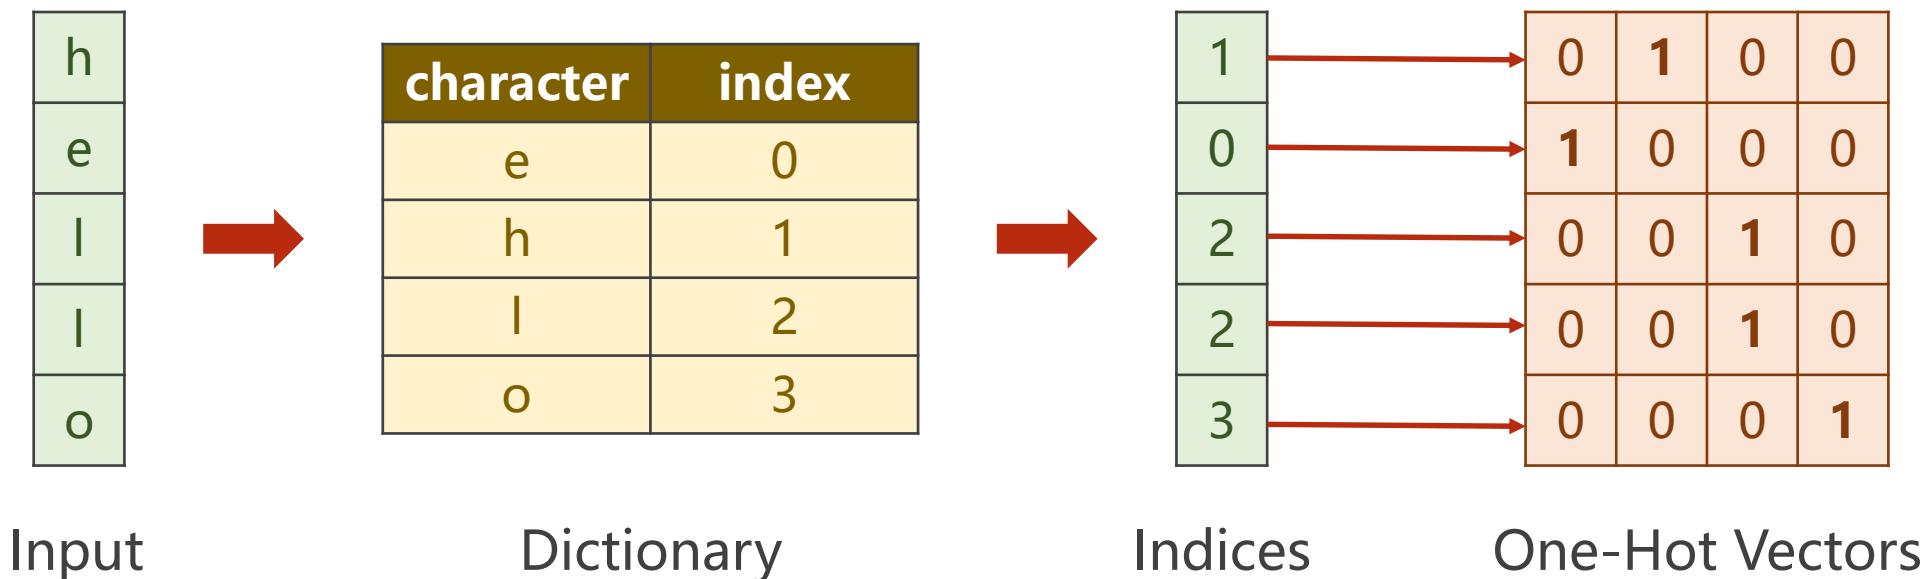
- The inputs of RNN Cell should be vectors of numbers.



*inputSize =?*

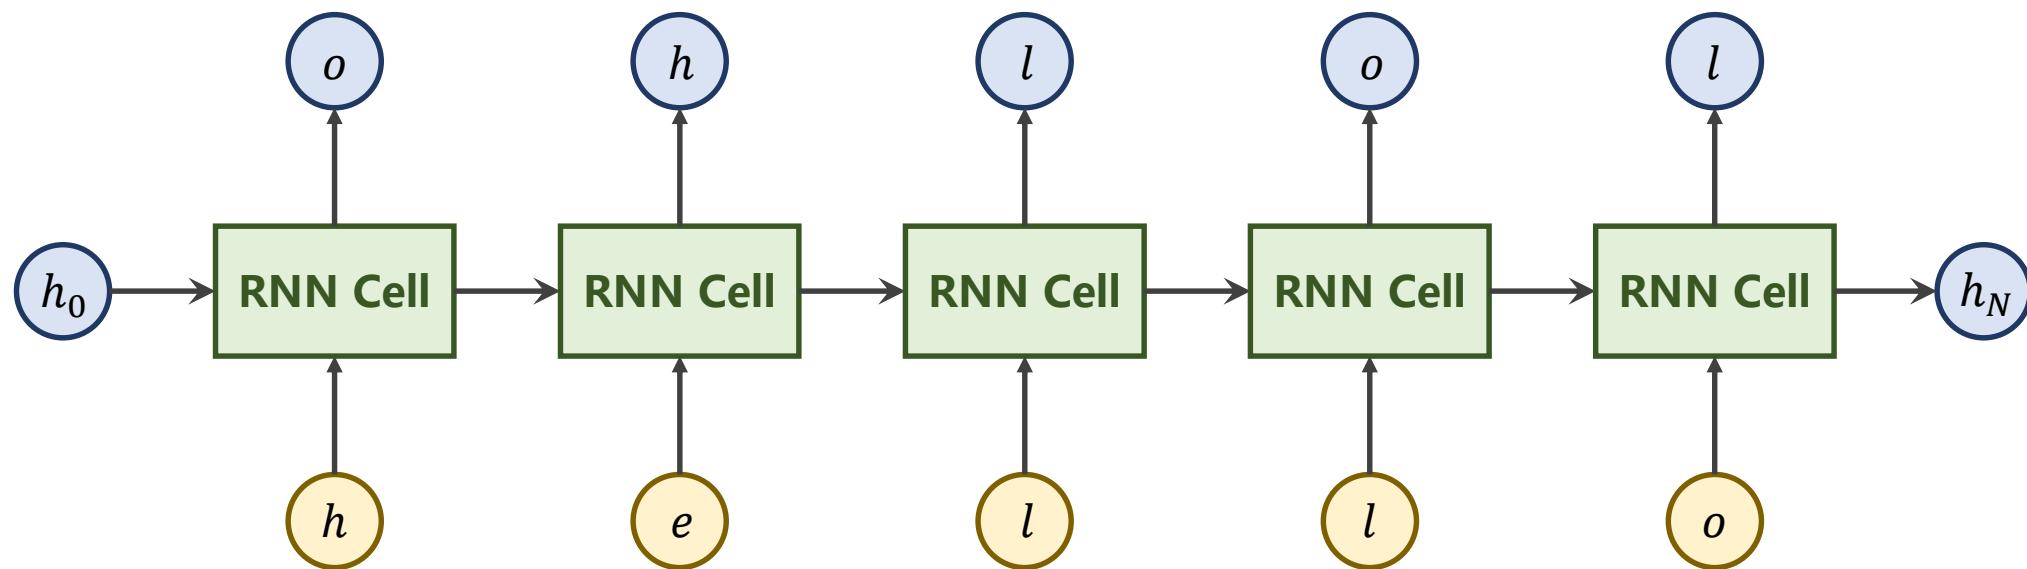
# Example 12-1: Using RNNCell

- The inputs of RNN Cell should be vectors of numbers.



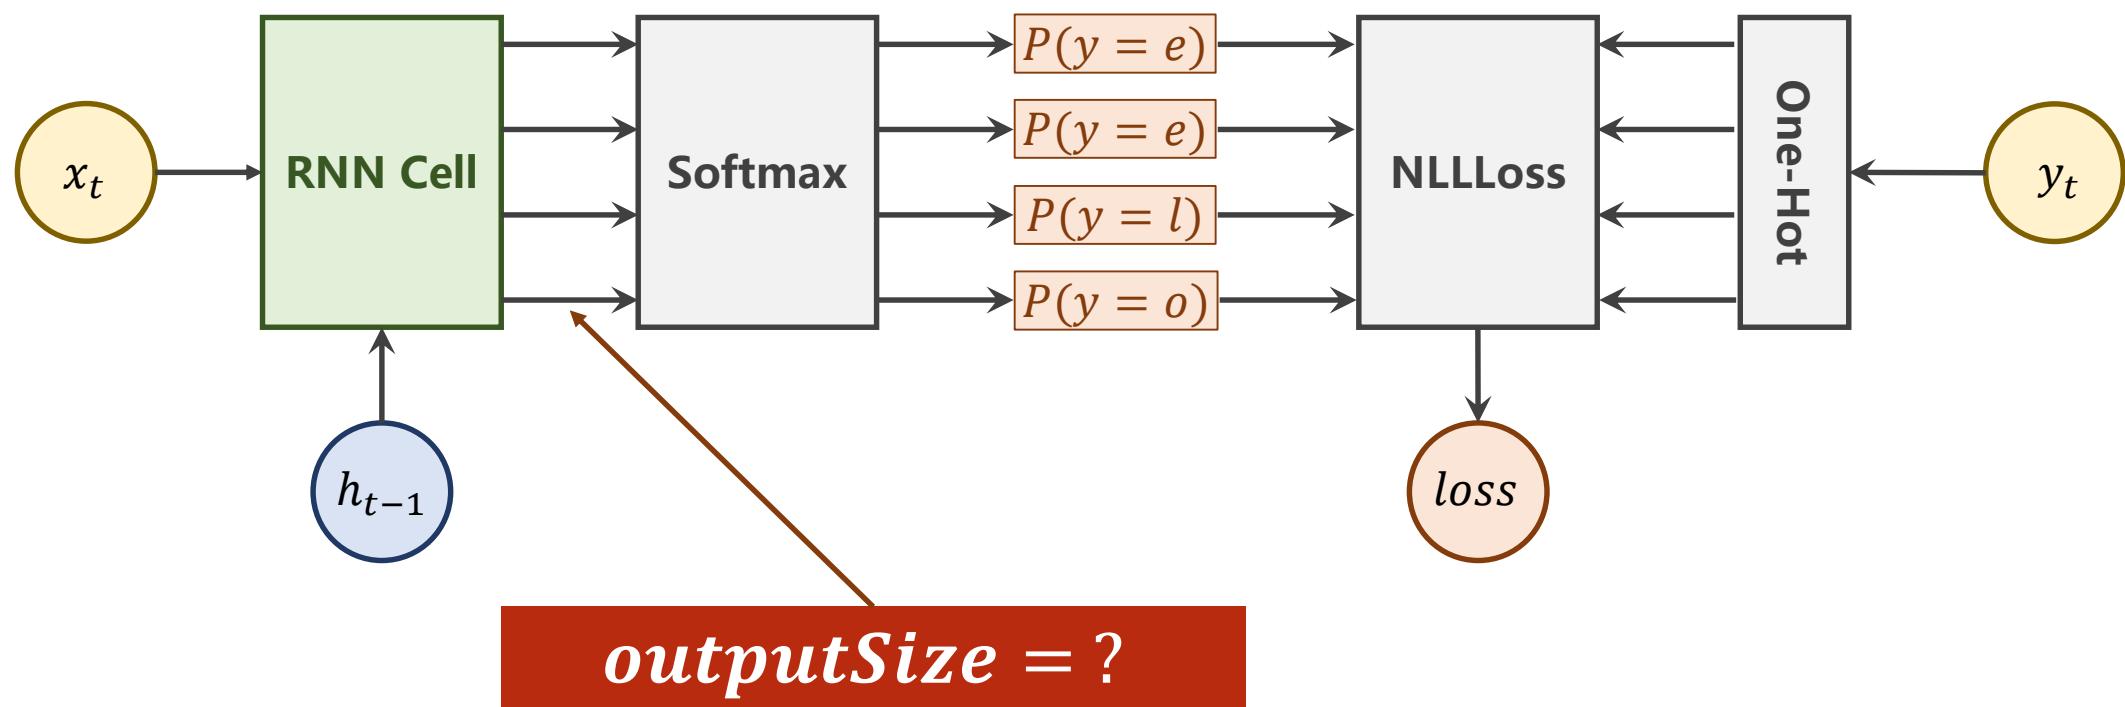
*inputSize = 4*

# Example 12-1: Using RNNCell



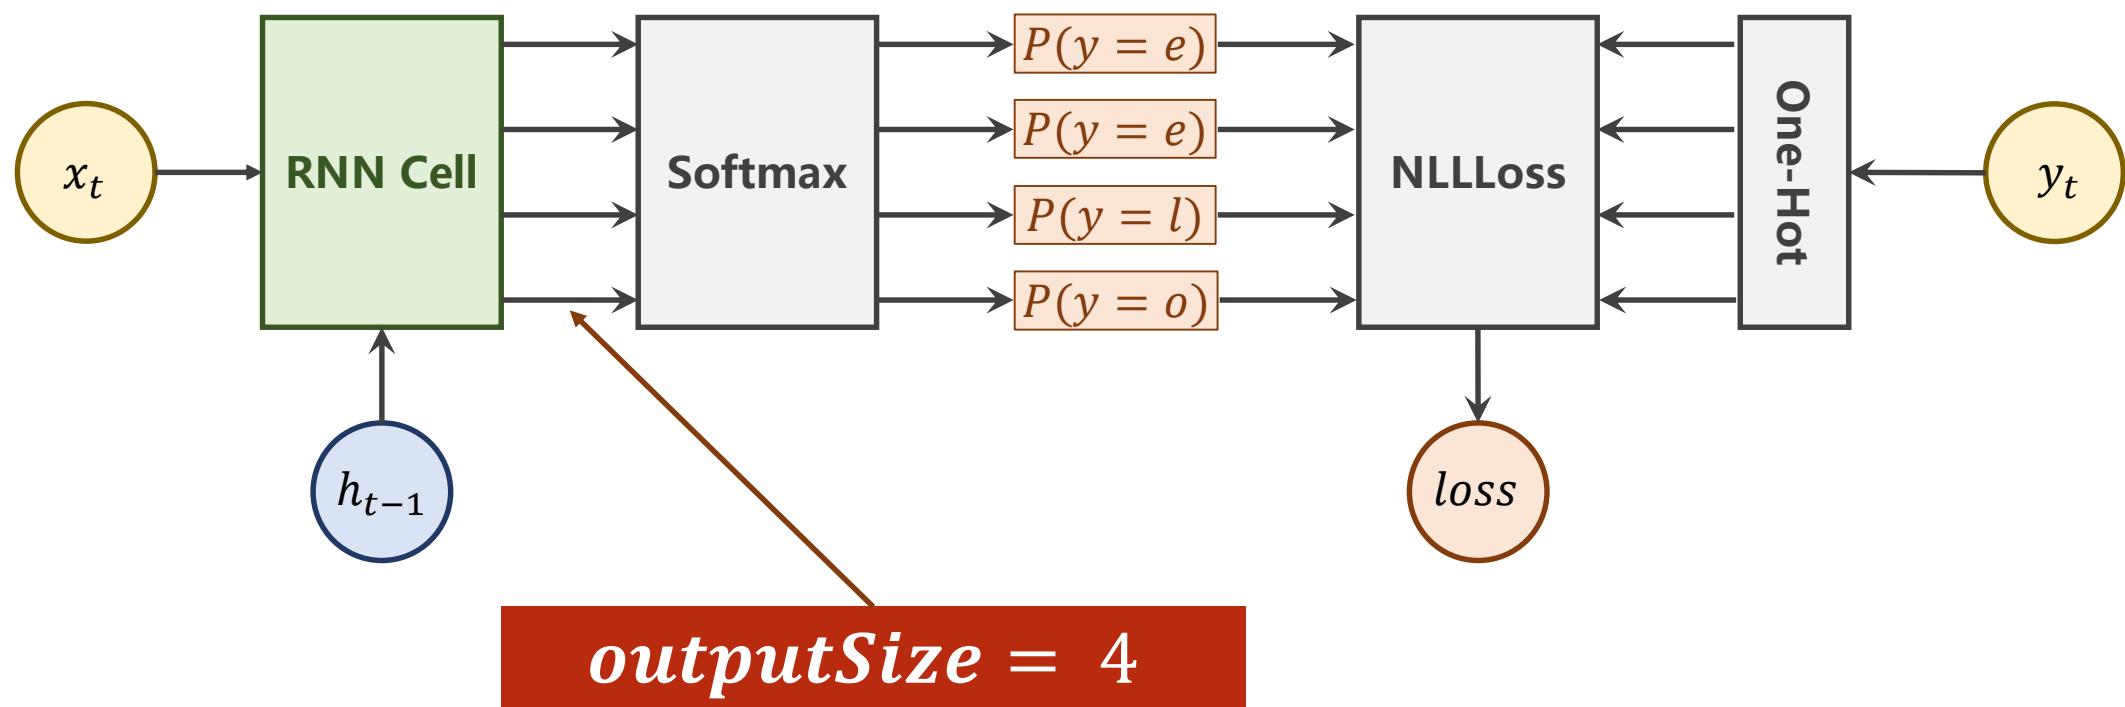
# Example 12-1: Using RNNCell

- The outputs of RNN Cell should be vectors of prediction.



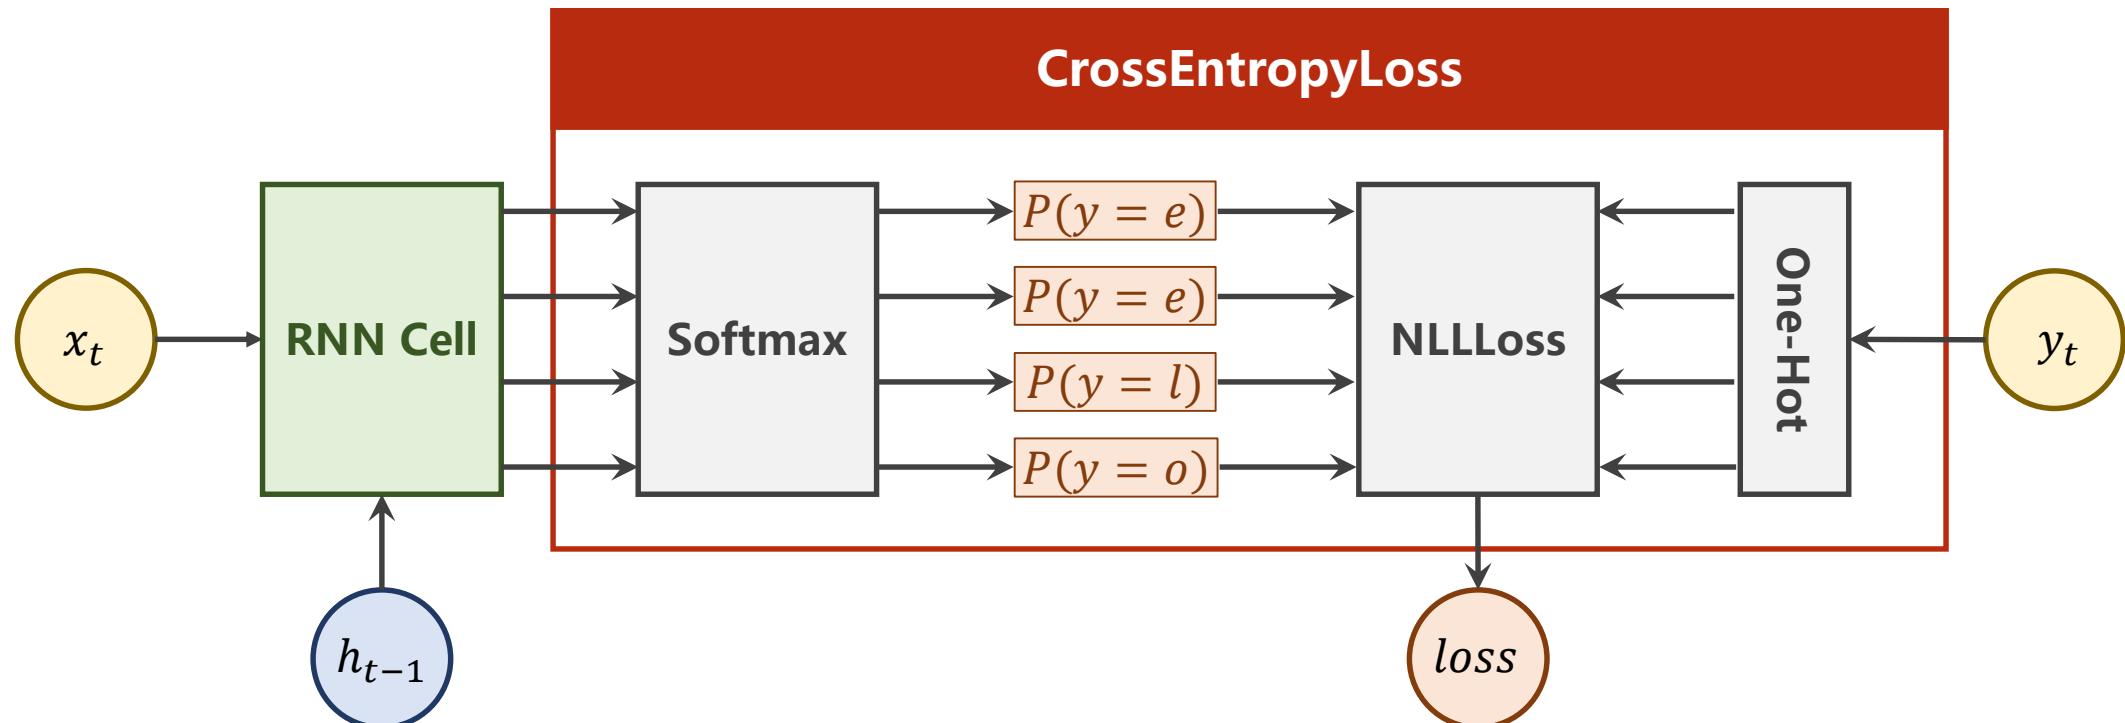
# Example 12-1: Using RNNCell

- The outputs of RNN Cell should be vectors of prediction.



# Example 12-1: Using RNNCell

- The outputs of RNN Cell should be vectors of prediction.



***outputSize = 4***

## Example 12-1: Code – Parameters

```
import torch  
  
input_size = 4  
hidden_size = 4  
batch_size = 1
```

# Example 12-1: Code – Prepare Data

```
idx2char = ['e', 'h', 'l', 'o']
x_data = [1, 0, 2, 2, 3]
y_data = [3, 1, 2, 3, 2]

one_hot_lookup = [[1, 0, 0, 0],
                  [0, 1, 0, 0],
                  [0, 0, 1, 0],
                  [0, 0, 0, 1]]
x_one_hot = [one_hot_lookup[x] for x in x_data]

inputs = torch.Tensor(x_one_hot).view(-1, batch_size, input_size)
labels = torch.LongTensor(y_data).view(-1, 1)
```

## The dictionary

# Example 12-1: Code – Prepare Data

```
idx2char = ['e', 'h', 'l', 'o']
x_data = [1, 0, 2, 2, 3]
y_data = [3, 1, 2, 3, 2]

one_hot_lookup = [[1, 0, 0, 0],
                  [0, 1, 0, 0],
                  [0, 0, 1, 0],
                  [0, 0, 0, 1]]
x_one_hot = [one_hot_lookup[x] for x in x_data]

inputs = torch.Tensor(x_one_hot).view(-1, batch_size, input_size)
labels = torch.LongTensor(y_data).view(-1, 1)
```

The input sequence is 'hello'

# Example 12-1: Code – Prepare Data

```
idx2char = ['e', 'h', 'l', 'o']
x_data = [1, 0, 2, 2, 3]
y_data = [3, 1, 2, 3, 2]

one_hot_lookup = [[1, 0, 0, 0],
                  [0, 1, 0, 0],
                  [0, 0, 1, 0],
                  [0, 0, 0, 1]]
x_one_hot = [one_hot_lookup[x] for x in x_data]

inputs = torch.Tensor(x_one_hot).view(-1, batch_size, input_size)
labels = torch.LongTensor(y_data).view(-1, 1)
```

The output sequence is 'ohlol'

# Example 12-1: Code – Prepare Data

```
idx2char = ['e', 'h', 'l', 'o']
x_data = [1, 0, 2, 2, 3]
y_data = [3, 1, 2, 3, 2]

one_hot_lookup = [[1, 0, 0, 0],
                  [0, 1, 0, 0],
                  [0, 0, 1, 0],
                  [0, 0, 0, 1]]
x_one_hot = [one_hot_lookup[x] for x in x_data]

inputs = torch.Tensor(x_one_hot).view(-1, batch_size, input_size)
labels = torch.LongTensor(y_data).view(-1, 1)
```

Convert indices into one-hot vector

# Example 12-1: Code – Prepare Data

```
idx2char = ['e', 'h', 'l', 'o']
x_data = [1, 0, 2, 2, 3]
y_data = [3, 1, 2, 3, 2]

one_hot_lookup = [[1, 0, 0, 0],
                  [0, 1, 0, 0],
                  [0, 0, 1, 0],
                  [0, 0, 0, 1]]
x_one_hot = [one_hot_lookup[x] for x in x_data]

inputs = torch.Tensor(x_one_hot).view(-1, batch_size, input_size)
labels = torch.LongTensor(y_data).view(-1, 1)
```

**Reshape the inputs to  
 $(seqLen, batchSize, inputSize)$**

# Example 12-1: Code – Prepare Data

```
idx2char = ['e', 'h', 'l', 'o']
x_data = [1, 0, 2, 2, 3]
y_data = [3, 1, 2, 3, 2]

one_hot_lookup = [[1, 0, 0, 0],
                  [0, 1, 0, 0],
                  [0, 0, 1, 0],
                  [0, 0, 0, 1]]
x_one_hot = [one_hot_lookup[x] for x in x_data]

inputs = torch.Tensor(x_one_hot).view(-1, batch_size, input_size)
labels = torch.LongTensor(y_data).view(-1, 1)
```

Reshape the labels to  
 $(seqLen, 1)$

# Example 12-1: Code – Design Model

```
class Model(torch.nn.Module):
    def __init__(self, input_size, hidden_size, batch_size):
        super(Model, self).__init__()
        self.batch_size = batch_size
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.rnncell1 = torch.nn.RNNCell(input_size=self.input_size,
                                         hidden_size=self.hidden_size)

    def forward(self, input, hidden):
        hidden = self.rnncell1(input, hidden)
        return hidden

    def init_hidden(self):
        return torch.zeros(self.batch_size, self.hidden_size)

net = Model(input_size, hidden_size, batch_size)
```

Initial the parameters

# Example 12-1: Code – Design Model

```
class Model(torch.nn.Module):
    def __init__(self, input_size, hidden_size, batch_size):
        super(Model, self).__init__()
        # self.num_layers = num_layers
        self.batch_size = batch_size
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.rnncell = torch.nn.RNNCell(input_size=self.input_size,
                                        hidden_size=self.hidden_size)

    def forward(self, input, hidden):
        hidden = self.rnncell(input, hidden)
        return hidden

    def init_hidden(self):
        return torch.zeros(self.batch_size, self.hidden_size)

net = Model(input_size, hidden_size, batch_size)
```

**Shape of inputs :**

*(batchSize, inputSize)*

**Shape of hidden:**

*(batchSize, hiddenSize)*

# Example 12-1: Code – Design Model

```
class Model(torch.nn.Module):
    def __init__(self, input_size, hidden_size, batch_size):
        super(Model, self).__init__()
        # self.num_layers = num_layers
        self.batch_size = batch_size
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.rnncell = torch.nn.RNNCell(input_size=self.input_size,
                                        hidden_size=self.hidden_size)

    def forward(self, input, hidden):
        hidden = self.rnncell(input, hidden)
        return hidden

    def init_hidden(self):
        return torch.zeros(self.batch_size, self.hidden_size)

net = Model(input_size, hidden_size, batch_size)
```

Provide initial hidden

# Example 12-1: Code – Loss and Optimizer

```
criterion = torch.nn.CrossEntropyLoss()  
optimizer = torch.optim.Adam(net.parameters(), lr=0.1)
```

# Example 12-1: Code – Training Cycle

```
for epoch in range(15):
    loss = 0
    optimizer.zero_grad()
    hidden = net.init_hidden()
    print('Predicted string: ', end='')
    for input, label in zip(inputs, labels):
        hidden = net(input, hidden)
        loss += criterion(hidden, label)
        _, idx = hidden.max(dim=1)
        print(idx2char[idx.item()], end='')
    loss.backward()
    optimizer.step()
    print(' Epoch [%d/15] loss=% .4f' % (epoch+1, loss.item()))
```

Training steps

# Example 12-1: Code – Training Cycle

```
for epoch in range(15):
    loss = 0
    optimizer.zero_grad()
    hidden = net.init_hidden()
    print('Predicted string: ', end='')
    for input, label in zip(inputs, labels):
        hidden = net(input, hidden)
        loss += criterion(hidden, label)
        _, idx = hidden.max(dim=1)
        print(idx2char[idx.item()], end=' ')
    loss.backward()
    optimizer.step()
    print(', Epoch [%d/15] loss=% .4f' % (epoch+1, loss.item()))
```

**Shape of *inputs*:**  
 $(seqLen, batchSize, inputSize)$

**Shape of *input*:**  
 $(batchSize, hiddenSize)$

# Example 12-1: Code – Training Cycle

```
for epoch in range(15):
    loss = 0
    optimizer.zero_grad()
    hidden = net.init_hidden()
    print('Predicted string: ', end='')
    for input, label in zip(inputs, labels):
        hidden = net(input, hidden)
        loss += criterion(hidden, label)
        _, idx = hidden.max(dim=1)
        print(idx2char[idx.item()], end=' ')
    loss.backward()
    optimizer.step()
    print(', Epoch [%d/15] loss=% .4f' % (epoch+1, loss.item()))
```

**Shape of *labels*:**  
 $(seqSize, 1)$

**Shape of *label*:**  
 $(1)$

# Example 12-1: Code – Training Cycle

```
for epoch in range(15):
    loss = 0
    optimizer.zero_grad()
    hidden = net.init_hidden()
    print('Predicted string: ', end='')
    for input, label in zip(inputs, labels):
        hidden = net(input, hidden)
        loss += criterion(hidden, label)
        _, idx = hidden.max(dim=1)
        print(idx2char[idx.item()], end='')
    loss.backward()
    optimizer.step()
    print(', Epoch [%d/15] loss=% .4f' % (epoch+1, loss.item()))
```

RNN Cell

# Example 12-1: Code – Training Cycle

```
for epoch in range(15):
    loss = 0
    optimizer.zero_grad()
    hidden = net.init_hidden()
    print('Predicted string: ', end='')
    for input, label in zip(inputs, labels):
        hidden = net(input, hidden)
        loss += criterion(hidden, label)
        _, idx = hidden.max(dim=1)
        print(idx2char[idx.item()], end='')
    loss.backward()
    optimizer.step()
    print(', Epoch [%d/15] loss=% .4f' % (epoch+1, loss.item()))
```

Output prediction

## Example 12-1: Code – Result

```
Predicted string: loeeh, Epoch [1/15] loss=8.0117
Predicted string: oooeh, Epoch [2/15] loss=7.2082
Predicted string: ooooh, Epoch [3/15] loss=6.6208
Predicted string: ooooh, Epoch [4/15] loss=6.1802
Predicted string: ooooh, Epoch [5/15] loss=5.8060
Predicted string: ooooh, Epoch [6/15] loss=5.4739
Predicted string: ooooh, Epoch [7/15] loss=5.1593
Predicted string: ooool, Epoch [8/15] loss=4.8593
Predicted string: ooool, Epoch [9/15] loss=4.5819
Predicted string: ohool, Epoch [10/15] loss=4.3287
Predicted string: ohlol, Epoch [11/15] loss=4.0909
Predicted string: ohlol, Epoch [12/15] loss=3.8573
Predicted string: ohlol, Epoch [13/15] loss=3.6246
Predicted string: ohlol, Epoch [14/15] loss=3.4007
Predicted string: ohlol, Epoch [15/15] loss=3.2005
```

# Example 12-2 Using RNN Module

```
criterion = torch.nn.CrossEntropyLoss()  
optimizer = torch.optim.Adam(net.parameters(), lr=0.05)  
  
for epoch in range(15):  
    optimizer.zero_grad()  
    outputs = net(inputs)  
    loss = criterion(outputs, labels)  
    loss.backward()  
    optimizer.step()  
  
    _, idx = outputs.max(dim=1)  
    idx = idx.data.numpy()  
    print('Predicted: ', ''.join([idx2char[x] for x in idx]), end=',')  
    print(', Epoch [%d/15] loss = %.3f' % (epoch + 1, loss.item()))
```

Training steps

# Example 12-2 Change Model

```
class Model(torch.nn.Module):
    def __init__(self, input_size, hidden_size, batch_size, num_layers=1):
        super(Model, self).__init__()
        self.num_layers = num_layers
        self.batch_size = batch_size
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.rnn = torch.nn.RNN(input_size=self.input_size,
                               hidden_size=self.hidden_size,
                               num_layers=num_layers)

    def forward(self, input):
        hidden = torch.zeros(self.num_layers,
                            self.batch_size,
                            self.hidden_size)
        out, _ = self.rnn(input, hidden)
        return out.view(-1, self.hidden_size)

net = Model(input_size, hidden_size, batch_size, num_layers)
```

A diagram illustrating the parameter assignment for the RNN layer. A curly brace on the left side of the code groups the parameters `input_size`, `hidden_size`, `num_layers`, and `batch_size`. An arrow points from this brace to a box containing their assigned values: `input_size = 4`, `hidden_size = 4`, `num_layers = 1`, `batch_size = 1`, and `seq_len = 5`.

input_size = 4
hidden_size = 4
num_layers = 1
batch_size = 1
seq_len = 5

# Example 12-2 Change Model

```
class Model(torch.nn.Module):
    def __init__(self, input_size, hidden_size, batch_size, num_layers=1):
        super(Model, self).__init__()
        self.num_layers = num_layers
        self.batch_size = batch_size
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.rnn = torch.nn.RNN(input_size=self.input_size,
                               hidden_size=self.hidden_size,
                               num_layers=num_layers)

    def forward(self, input):
        hidden = torch.zeros(self.num_layers,
                            self.batch_size,
                            self.hidden_size)
        out, _ = self.rnn(input, hidden)
        return out.view(-1, self.hidden_size)

net = Model(input_size, hidden_size, batch_size, num_layers)
```

**input\_size** = 4  
**hidden\_size** = 4  
**num\_layers** = 1  
batch\_size = 1  
seq\_len = 5

# Example 12-2 Change Model

```
class Model(torch.nn.Module):
    def __init__(self, input_size, hidden_size, batch_size, num_layers=1):
        super(Model, self).__init__()
        self.num_layers = num_layers
        self.batch_size = batch_size
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.rnn = torch.nn.RNN(input_size=self.input_size,
                               hidden_size=self.hidden_size,
                               num_layers=num_layers)
```

```
def forward(self, input):
    hidden = torch.zeros(self.num_layers,
                        self.batch_size,
                        self.hidden_size)
    out, _ = self.rnn(input, hidden)
    return out.view(-1, self.hidden_size)
```

```
net = Model(input_size, hidden_size, batch_size, num_layers)
```

```
input_size = 4
hidden_size = 4
num_layers = 1
batch_size = 1
seq_len = 5
```

**Shape of *hidden*:**  
 $(numLayers, batchSize, hiddenSize)$

# Example 12-2 Change Model

```
class Model(torch.nn.Module):
    def __init__(self, input_size, hidden_size, batch_size, num_layers=1):
        super(Model, self).__init__()
        self.num_layers = num_layers
        self.batch_size = batch_size
        self.input_size = input_size
        self.hidden_size = hidden_size
        self.rnn = torch.nn.RNN(input_size=self.input_size,
                               hidden_size=self.hidden_size,
                               num_layers=num_layers)

    def forward(self, input):
        hidden = torch.zeros(self.num_layers,
                            self.batch_size,
                            self.hidden_size)
        out, _ = self.rnn(input, hidden)
        return out.view(-1, self.hidden_size)

net = Model(input_size, hidden_size, batch_size, num_layers)
```

**Reshape out to:**

$(seqLen \times batchSize, hiddenSize)$

## Example 12-2 Change Data

```
idx2char = [ 'e', 'h', 'l', 'o' ]
```

```
x_data = [1, 0, 2, 2, 3]
```

```
y_data = [3, 1, 2, 3, 2]
```

```
one_hot_lookup = [[1, 0, 0, 0],  
                  [0, 1, 0, 0],  
                  [0, 0, 1, 0],  
                  [0, 0, 0, 1]]
```

```
x_one_hot = [one_hot_lookup[x] for x in x_data]
```

```
inputs = torch.Tensor(x_one_hot).view(seq_len, batch_size, input_size)
```

```
labels = torch.LongTensor(y_data)
```

```
input_size = 4  
hidden_size = 4  
num_layers = 1  
batch_size = 1  
seq_len = 5
```

**Shape of *inputs*:**  
 $(seqLen, batchSize, hiddenSize)$

## Example 12-2 Change Data

```
idx2char = [ 'e' , 'h' , 'l' , 'o' ]  
x_data = [1, 0, 2, 2, 3]  
y_data = [3, 1, 2, 3, 2]
```

```
one_hot_lookup = [[1, 0, 0, 0],  
                  [0, 1, 0, 0],  
                  [0, 0, 1, 0],  
                  [0, 0, 0, 1]]
```

```
x_one_hot = [one_hot_lookup[x] for x in x_data]
```

```
inputs = torch.Tensor(x_one_hot).view(seq_len, batch_size, input_size)  
labels = torch.LongTensor(y_data)
```

```
input_size = 4  
hidden_size = 4  
num_layers = 1  
batch_size = 1  
seq_len = 5
```

**Shape of labels:**

$(seqLen \times batchSize, 1)$

# Example 12-2 Result

```
Predicted: Ihhhh, Epoch [1/15] loss = 1.325
Predicted: lohhh, Epoch [2/15] loss = 1.190
Predicted: lolol, Epoch [3/15] loss = 1.090
Predicted: lolol, Epoch [4/15] loss = 1.017
Predicted: oolol, Epoch [5/15] loss = 0.958
Predicted: oolol, Epoch [6/15] loss = 0.904
Predicted: oolol, Epoch [7/15] loss = 0.852
Predicted: oolol, Epoch [8/15] loss = 0.804
Predicted: oolol, Epoch [9/15] loss = 0.759
Predicted: ohlol, Epoch [10/15] loss = 0.716
Predicted: ohlol, Epoch [11/15] loss = 0.678
Predicted: ohlol, Epoch [12/15] loss = 0.643
Predicted: ohlol, Epoch [13/15] loss = 0.613
Predicted: ohlol, Epoch [14/15] loss = 0.588
Predicted: ohlol, Epoch [15/15] loss = 0.568
```

# Associate a vector with a word/character

- One-hot encoding of words and characters
  - The one-hot vectors are **high-dimension**. 独热向量维度太高
  - The one-hot vectors are **sparse**. 过于稀疏
  - The one-hot vectors are **hardcoded**. 硬编码

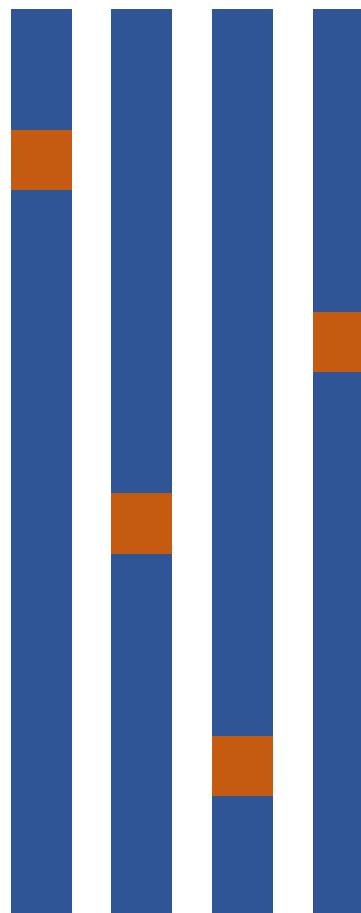
# Associate a vector with a word/character

- One-hot encoding of words and characters
  - The one-hot vectors are **high-dimension**.
  - The one-hot vectors are **sparse**.
  - The one-hot vectors are **hardcoded**.
- Do we have a way to associate a vector with a word/character with following specification:
  - **Lower-dimension**
  - **Dense**
  - **Learned from data**
- A popular and powerful way is called **EMBEDDING**.

# Associate a vector with a word/character

- One-hot encoding of words and characters
  - The one-hot vectors are high-dimension.
  - The one-hot vectors are sparse.
  - The one-hot vectors are hardcoded.
- Do we have a way to associate a vector with a word/character with following specification:
  - Lower-dimension
  - Dense
  - Learned from data
- A popular and powerful way is called **EMBEDDING**.

# One-hot vs Embedding



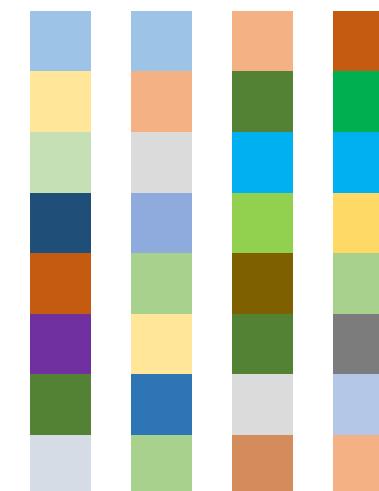
## One-hot vectors:

- High-dimension
- Sparse
- Hardcoded

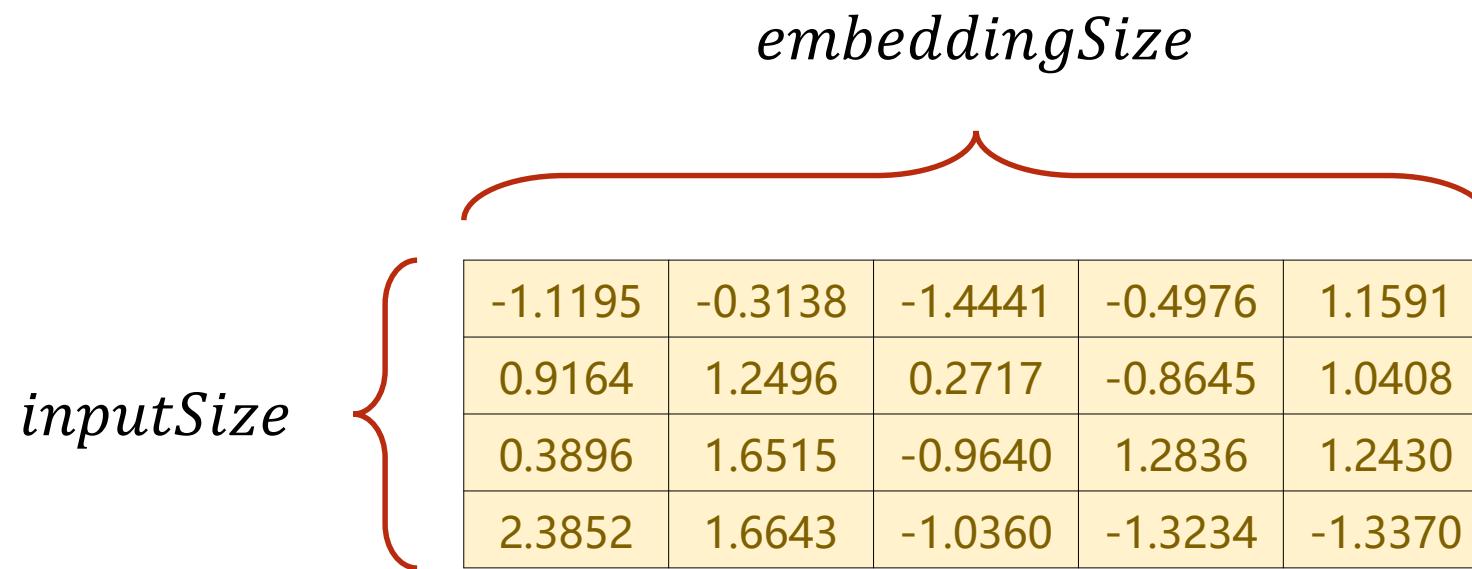
将高维的稀疏数据降维到低维的稠密数据

## Embedding vectors:

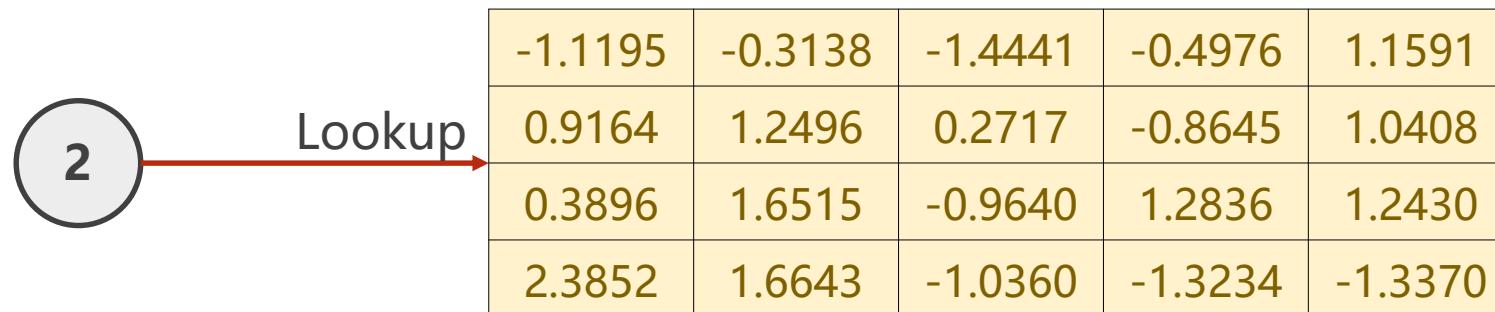
- Lower-dimension
- Dense
- Learned from data



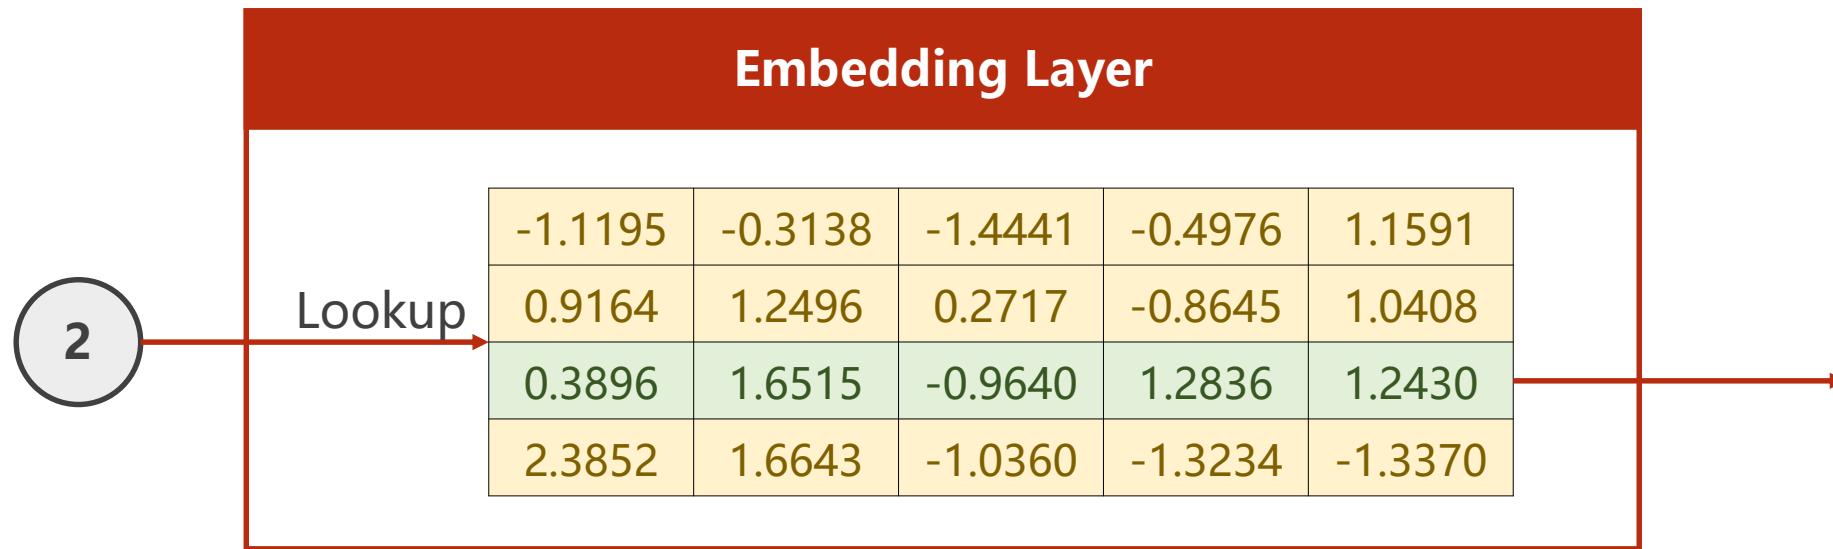
# Embedding in PyTorch



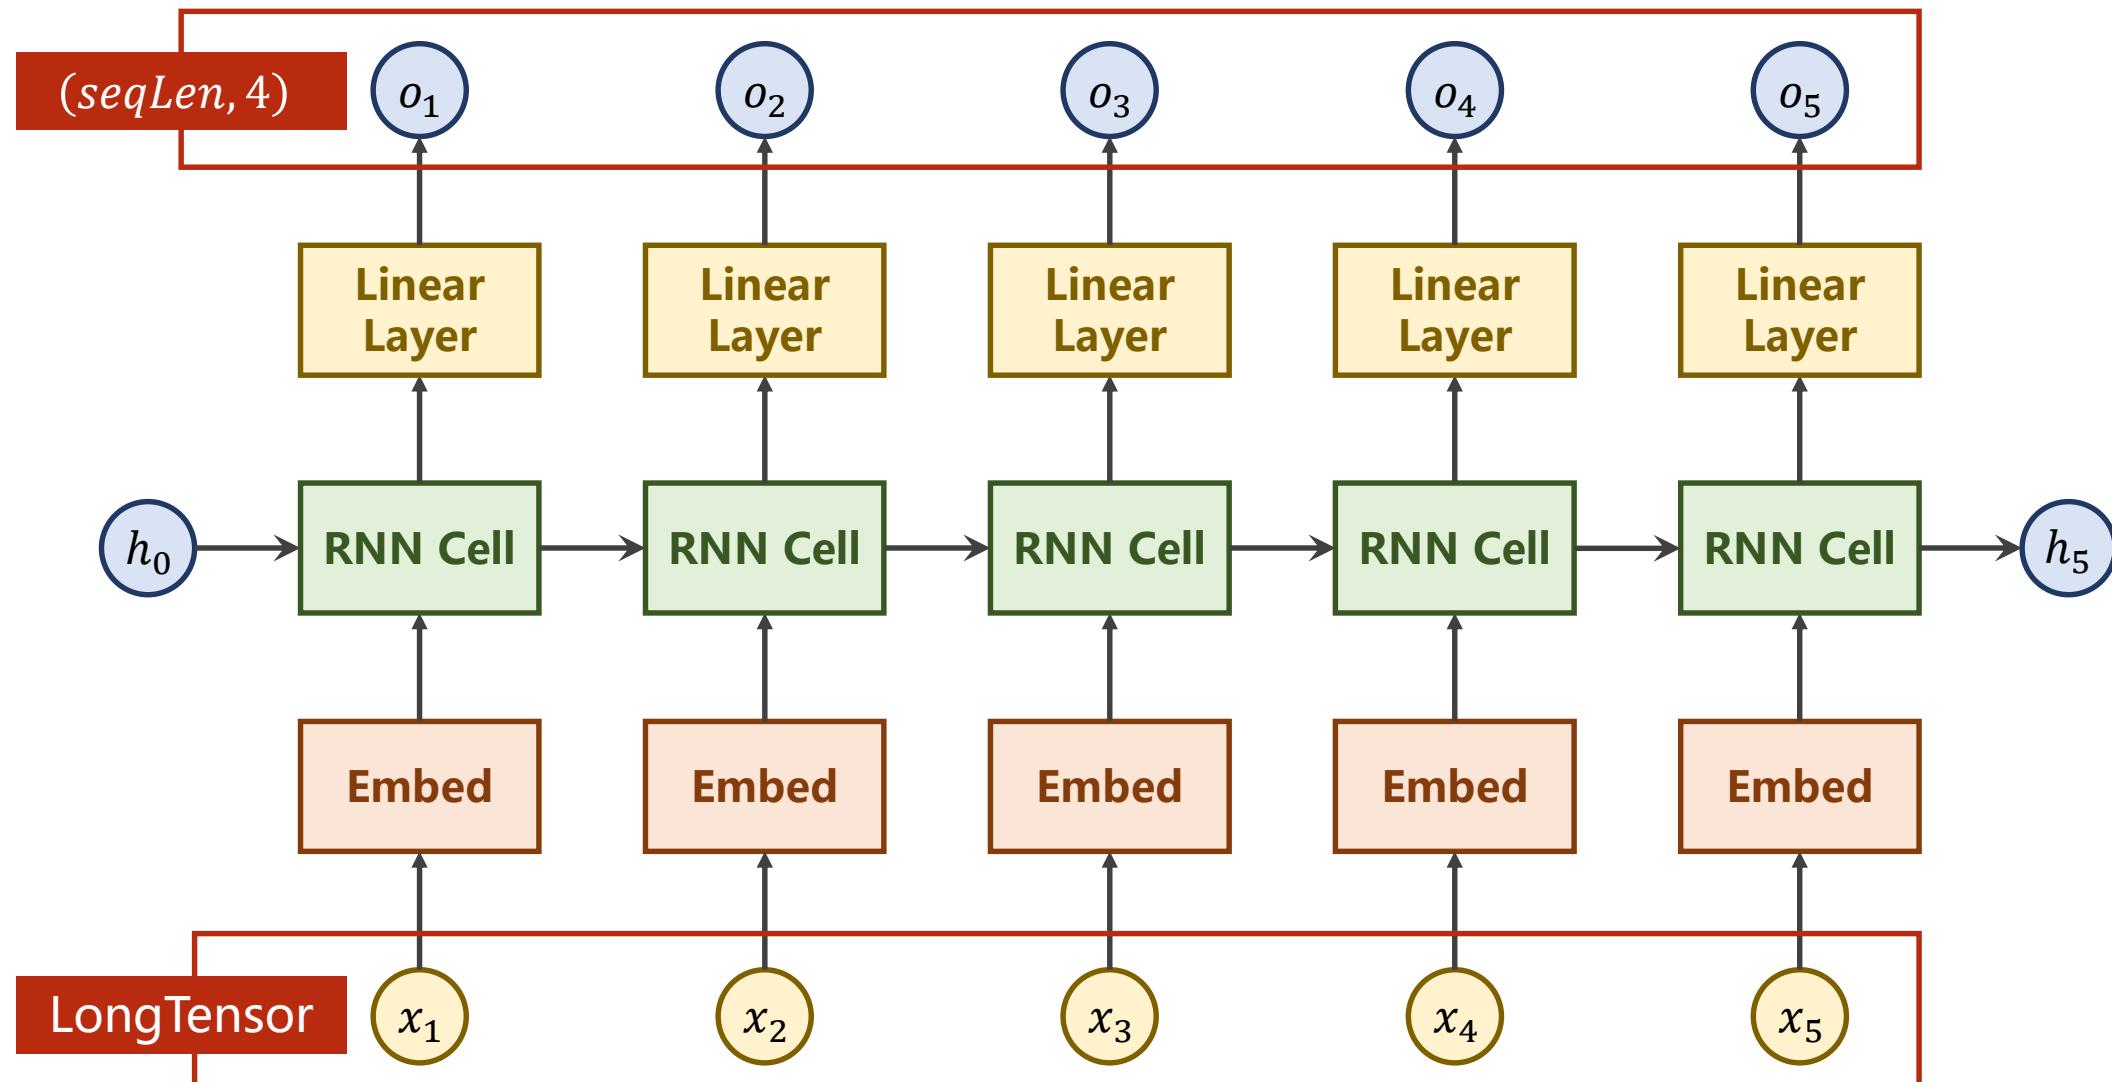
# Embedding in PyTorch



# Embedding in PyTorch



# Example 12-3 Using embedding and linear layer



# Example 12-3 Using embedding and linear layer

```
class torch.nn.Embedding(num_embeddings, embedding_dim, padding_idx=None, max_norm=None,  
norm_type=2, scale_grad_by_freq=False, sparse=False, _weight=None) [source]
```

A simple lookup table that stores embeddings of a fixed dictionary and size.

This module is often used to store word embeddings and retrieve them using indices. The input to the module is a list of indices, and the output is the corresponding word embeddings.

Parameters:

- num\_embeddings (*int*) – size of the dictionary of embeddings
- embedding\_dim (*int*) – the size of each embedding vector
- padding\_idx (*int, optional*) – If given, pads the output with the embedding vector at

Shape:

- Input: LongTensor of arbitrary shape containing the indices to extract
- Output:  $(*, \text{embedding\_dim})$ , where  $*$  is the input shape

- sparse (*bool, optional*) – if `True`, gradient w.r.t. weight matrix will be a sparse tensor. See Notes for more details regarding sparse gradients.

# Example 12-3 Using embedding and linear layer

```
class torch.nn.Linear(in_features, out_features, bias=True) [source]
```

Applies a linear transformation to the incoming data:  $y = Ax + b$

Parameters:

- **in\_features** – size of each input sample
- **out\_features** – size of each output sample
- **bias** – If set to False, the layer will not learn an additive bias. Default: `True`

Shape:

- Input:  $(N, *, in\_features)$  where \* means any number of additional dimensions
- Output:  $(N, *, out\_features)$  where all but the last dimension are the same shape as the input.

# Example 12-3 Using embedding and linear layer

```
class torch.nn.CrossEntropyLoss (weight=None, size_average=True, ignore_index=-100, reduce=True)
\[source\]
```

This criterion combines `nn.LogSoftmax()` and `nn.NLLLoss()` in one single class.

## Shape:

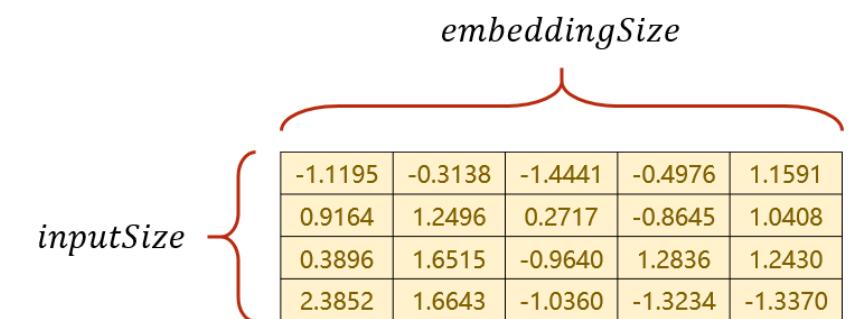
- Input:  $(N, C)$  where  $C = \text{number of classes}$ , or  
 $(N, C, d_1, d_2, \dots, d_K)$  with  $K \geq 2$  in the case of K-dimensional loss.
- Target:  $(N)$  where each value is  $0 \leq \text{targets}[i] \leq C - 1$ , or  
 $(N, d_1, d_2, \dots, d_K)$  with  $K \geq 2$  in the case of K-dimensional loss.
- Output: scalar. If `reduce` is `False`, then the same size  
as the target:  $(N)$ , or  $(N, d_1, d_2, \dots, d_K)$  with  $K \geq 2$  in the case of K-dimensional loss.

# Example 12-3 Using embedding and linear layer

```
class Model(torch.nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.emb = torch.nn.Embedding(input_size, embedding_size)
        self.rnn = torch.nn.RNN(input_size=embedding_size,
                               hidden_size=hidden_size,
                               num_layers=num_layers,
                               batch_first=True)
        self.fc = torch.nn.Linear(hidden_size, num_class)

    def forward(self, x):
        hidden = torch.zeros(num_layers, x.size(0), hidden_size)
        x = self.emb(x) # (batch, seqLen, embeddingSize)
        x, _ = self.rnn(x, hidden)
        x = self.fc(x)
        return x.view(-1, num_class)
```

**Lookup matrix of Embedding:**  
*(inputSize, embeddingSize)*



embeddingSize				
-1.1195	-0.3138	-1.4441	-0.4976	1.1591
0.9164	1.2496	0.2717	-0.8645	1.0408
0.3896	1.6515	-0.9640	1.2836	1.2430
2.3852	1.6643	-1.0360	-1.3234	-1.3370

# Example 12-3 Using embedding and linear layer

```
class Model(torch.nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.emb = torch.nn.Embedding(input_size, embedding_size)
        self.rnn = torch.nn.RNN(input_size=embedding_size,
                               hidden_size=hidden_size,
                               num_layers=num_layers,
                               batch_first=True)
        self.fc = torch.nn.Linear(hidden_size, num_class)

    def forward(self, x):
        hidden = torch.zeros(num_layers, x.size(0), hidden_size)
        x = self.emb(x) # (batch, seqLen, embeddingSize)
        x, _ = self.rnn(x, hidden)
        x = self.fc(x)
        return x.view(-1, num_class)
```

**Input should be LongTensor:**

*(batchSize, seqLen)*

**Output with shape:**

*(batchSize, seqLen, embeddingSize)*

**Notice: batch FIRST**

# Example 12-3 Using embedding and linear layer

```
class Model(torch.nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.emb = torch.nn.Embedding(input_size, embedding_size)
        self.rnn = torch.nn.RNN(input_size=embedding_size,
                               hidden_size=hidden_size,
                               num_layers=num_layers,
                               batch_first=True)
        self.fc = torch.nn.Linear(hidden_size, num_class)

    def forward(self, x):
        hidden = torch.zeros(num_layers, x.size(0), hidden_size)
        x = self.emb(x) # (batch, seqLen, embeddingSize)
        x, _ = self.rnn(x, hidden)
        x = self.fc(x)
        return x.view(-1, num_class)
```

**Input of RNN:**

*(batchSize, seqLen, embeddingSize)*

**Output of RNN:**

*(batchSize, seqLen, hiddenSize)*

# Example 12-3 Using embedding and linear layer

```
class Model(torch.nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.emb = torch.nn.Embedding(input_size, embedding_size)
        self.rnn = torch.nn.RNN(input_size=embedding_size,
                               hidden_size=hidden_size,
                               num_layers=num_layers,
                               batch_first=True)
        self.fc = torch.nn.Linear(hidden_size, num_class)

    def forward(self, x):
        hidden = torch.zeros(num_layers, x.size(0), hidden_size)
        x = self.emb(x) # (batch, seqLen, embeddingSize)
        x, _ = self.rnn(x, hidden)
        x = self.fc(x)
        return x.view(-1, num_class)
```

## Shape:

- Input:  $(N, *, in\_features)$  where  $N$  is the batch size.
- Output:  $(N, *, out\_features)$  where  $*$  depends on the input.

## Input of FC Layer:

$(batchSize, seqLen, hiddenSize)$

## Output of FC Layer:

$(batchSize, seqLen, numClass)$

# Example 12-3 Using embedding and linear layer

```
class Model(torch.nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.emb = torch.nn.Embedding(input_size, embedding_size)
        self.rnn = torch.nn.RNN(input_size=embedding_size,
                               hidden_size=hidden_size,
                               num_layers=num_layers,
                               batch_first=True)
        self.fc = torch.nn.Linear(hidden_size, num_class)

    def forward(self, x):
        hidden = torch.zeros(num_layers, x.size(0), hidden_size)
        x = self.emb(x) # (batch, seqLen, embeddingSize)
        x, _ = self.rnn(x, hidden)
        x = self.fc(x)
        return x.view(-1, num_class)
```

Shape:

- Input:  $(N, C)$  where  $N = \text{batchSize}$ ,  $C = \text{seqLen}$   
 $(N, C, d_1, d_2, \dots)$
- Target:  $(N)$  where  $N = \text{batchSize} \times \text{seqLen}$   
 $(N, d_1, d_2, \dots)$

Reshape result to use Cross Entropy

Loss:

$(\text{batchSize} \times \text{seqLen}, \text{numClass})$

## Example 12-3 Using embedding and linear layer

```
# parameters
num_class = 4
input_size = 4
hidden_size = 8
embedding_size = 10
num_layers = 2
batch_size = 1
seq_len = 5
```

# Example 12-3 Using embedding and linear layer

```
idx2char = ['e', 'h', 'l', 'o']
x_data = [[1, 0, 2, 2, 3]] # (batch, seq_len)
y_data = [3, 1, 2, 3, 2] # (batch * seq_len)

inputs = torch.LongTensor(x_data)
labels = torch.LongTensor(y_data)
```

**Input should be LongTensor:**

$(batchSize, seqLen)$

**Target should be LongTensor:**

$(batchSize \times seqLen)$

```
# parameters
num_class = 4
input_size = 4
hidden_size = 8
embedding_size = 10
num_layers = 2
batch_size = 1
seq_len = 5
```

# Example 12-3 Using embedding and linear layer

```
net = Model()  
  
criterion = torch.nn.CrossEntropyLoss()  
optimizer = torch.optim.Adam(net.parameters(), lr=0.05)
```

# Example 12-3 Using embedding and linear layer

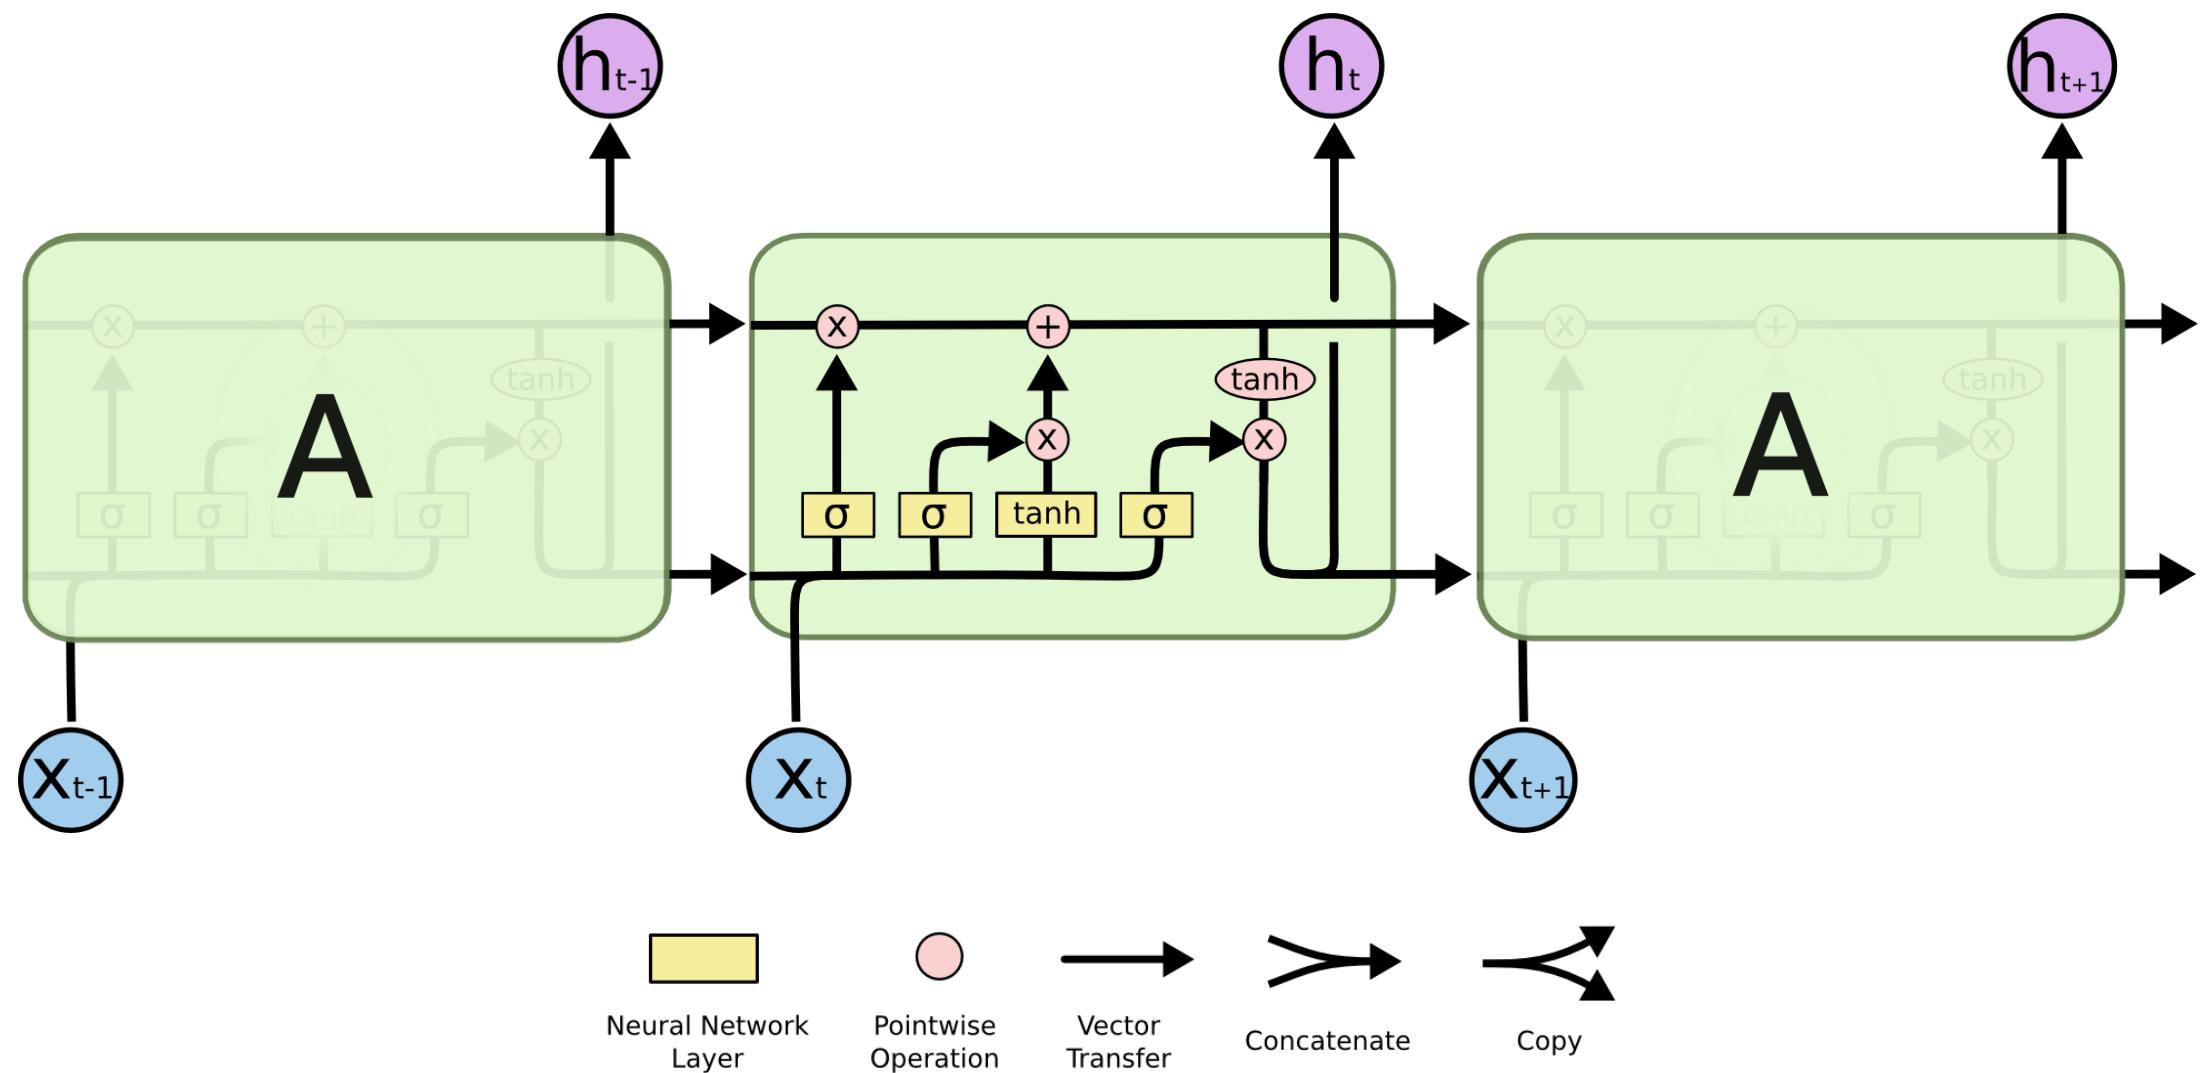
```
for epoch in range(15):
    optimizer.zero_grad()
    outputs = net(inputs)
    loss = criterion(outputs, labels)
    loss.backward()
    optimizer.step()

    _, idx = outputs.max(dim=1)
    idx = idx.data.numpy()
    print('Predicted: ', ''.join([idx2char[x] for x in idx]), end=' ')
    print(', Epoch [%d/15] loss = %.3f' % (epoch + 1, loss.item()))
```

# Example 12-3 Using embedding and linear layer

```
Predicted: IIII, Epoch [1/15] loss = 1.412
Predicted: IIII, Epoch [2/15] loss = 1.038
Predicted: IhIII, Epoch [3/15] loss = 0.799
Predicted: ohlol, Epoch [4/15] loss = 0.652
Predicted: ohloo, Epoch [5/15] loss = 0.548
Predicted: ohlol, Epoch [6/15] loss = 0.459
Predicted: ohlol, Epoch [7/15] loss = 0.381
Predicted: ohlol, Epoch [8/15] loss = 0.313
Predicted: ohlol, Epoch [9/15] loss = 0.248
Predicted: ohlol, Epoch [10/15] loss = 0.188
Predicted: ohlol, Epoch [11/15] loss = 0.136
Predicted: ohlol, Epoch [12/15] loss = 0.098
Predicted: ohlol, Epoch [13/15] loss = 0.072
Predicted: ohlol, Epoch [14/15] loss = 0.055
Predicted: ohlol, Epoch [15/15] loss = 0.043
```

# Exercise 12 – 1 Using LSTM



# Exercise 12 – 1 Using LSTM

`class torch.nn.LSTM(*args, **kwargs)` [\[source\]](#)

Applies a multi-layer long short-term memory (LSTM) RNN to an input sequence.

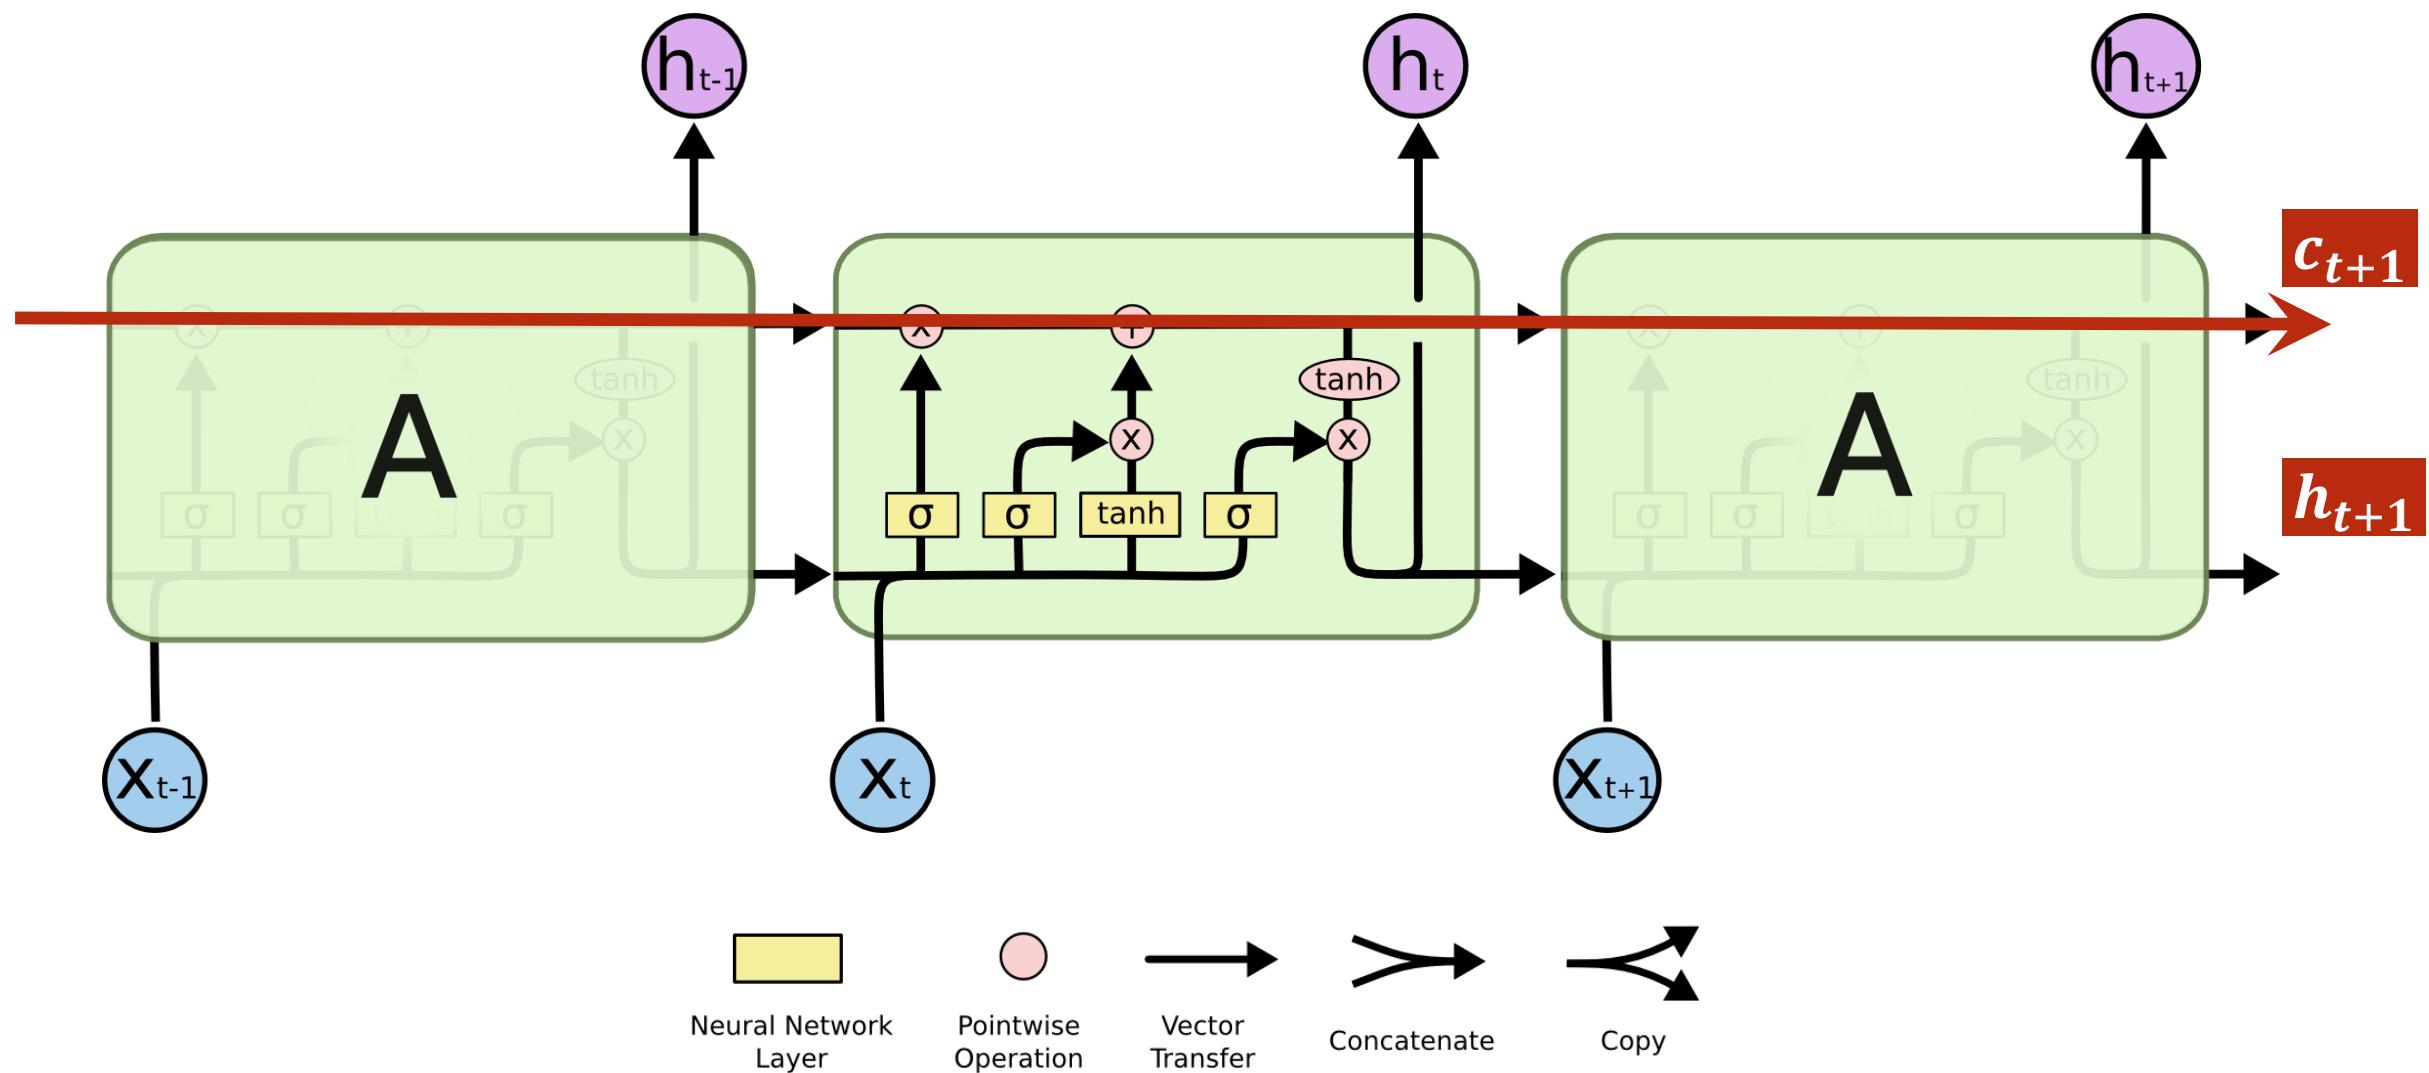
For each element in the input sequence, each layer computes the following function:

$$\begin{aligned} i_t &= \sigma(W_{ii}x_t + b_{ii} + W_{hi}h_{(t-1)} + b_{hi}) \\ f_t &= \sigma(W_{if}x_t + b_{if} + W_{hf}h_{(t-1)} + b_{hf}) \\ g_t &= \tanh(W_{ig}x_t + b_{ig} + W_{hg}h_{(t-1)} + b_{hg}) \\ o_t &= \sigma(W_{io}x_t + b_{io} + W_{ho}h_{(t-1)} + b_{ho}) \\ c_t &= f_t c_{(t-1)} + i_t g_t \\ h_t &= o_t \tanh(c_t) \end{aligned}$$

where  $h_t$  is the hidden state at time  $t$ ,  $c_t$  is the cell state at time  $t$ ,  $x_t$  is the input at time  $t$ ,  $h_{(t-1)}$  is the hidden state of the previous layer at time  $t-1$  or the initial hidden state at time 0, and  $i_t$ ,  $f_t$ ,  $g_t$ ,  $o_t$  are the input, forget, cell, and output gates, respectively.  $\sigma$  is the sigmoid function.

<https://pytorch.org/docs/stable/nn.html#lstm>

# Exercise 12 – 1 Using LSTM



# Exercise 12 – 1 Using LSTM

## Inputs: input, (h\_0, c\_0)

- **input** of shape  $(seq\_len, batch, input\_size)$ : tensor containing the features of the input sequence. The input can also be a packed variable length sequence. See `torch.nn.utils.rnn.pack_padded_sequence()` or `torch.nn.utils.rnn.pack_sequence()` for details.
- **h\_0** of shape  $(num\_layers * num\_directions, batch, hidden\_size)$ : tensor containing the initial hidden state for each element in the batch.
- **c\_0** of shape  $(num\_layers * num\_directions, batch, hidden\_size)$ : tensor containing the initial cell state for each element in the batch.

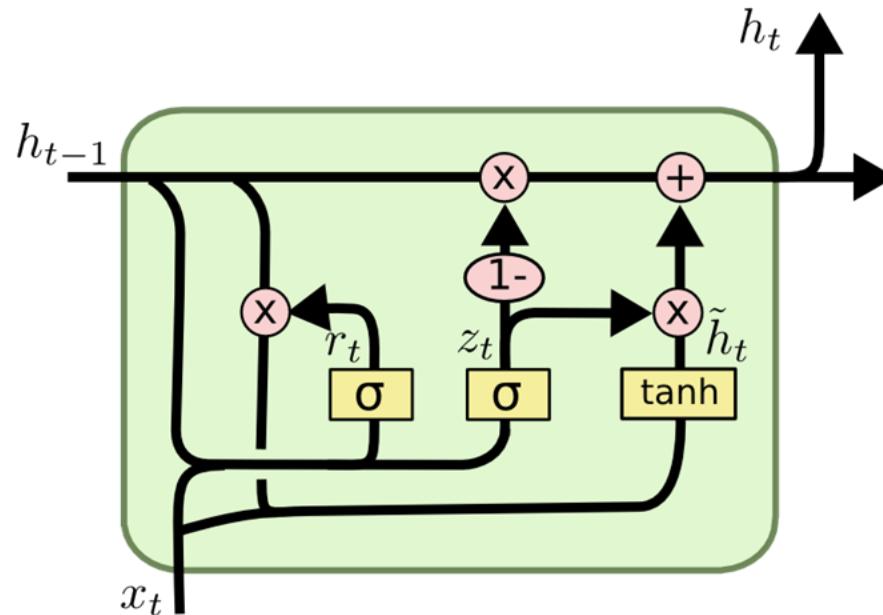
If  $(h_0, c_0)$  is not provided, both **h\_0** and **c\_0** default to zero.

## Outputs: output, (h\_n, c\_n)

- **output** of shape  $(seq\_len, batch, hidden\_size * num\_directions)$ : tensor containing the output features ( $h_t$ ) from the last layer of the LSTM, for each t. If a `torch.nn.utils.rnn.PackedSequence` has been given as the input, the output will also be a packed sequence.
- **h\_n** of shape  $(num\_layers * num\_directions, batch, hidden\_size)$ : tensor containing the hidden state for  $t = seq\_len$
- **c\_n** ( $num\_layers * num\_directions, batch, hidden\_size$ ): tensor containing the cell state for  $t = seq\_len$

<https://pytorch.org/docs/stable/nn.html#lstm>

# Exercise 12 – 2 Using GRU

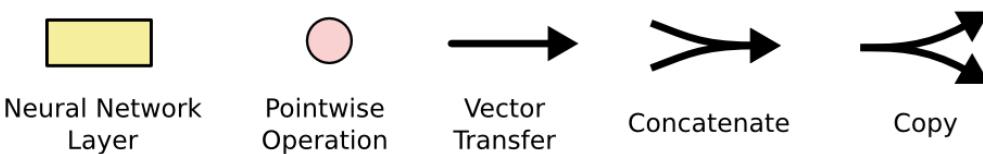


$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$



# Exercise 12 – 2 Using GRU

`class torch.nn.GRU(*args, **kwargs)` [source]

Applies a multi-layer gated recurrent unit (GRU) RNN to an input sequence.

For each element in the input sequence, each layer computes the following function:

$$r_t = \sigma(W_{ir}x_t + b_{ir} + W_{hr}h_{(t-1)} + b_{hr})$$

$$z_t = \sigma(W_{iz}x_t + b_{iz} + W_{hz}h_{(t-1)} + b_{hz})$$

$$n_t = \tanh(W_{in}x_t + b_{in} + r_t(W_{hn}h_{(t-1)} + b_{hn}))$$

$$h_t = (1 - z_t)n_t + z_t h_{(t-1)}$$

where  $h_t$  is the hidden state at time  $t$ ,  $x_t$  is the input at time  $t$ ,  $h_{(t-1)}$  is the hidden state of the previous layer at time  $t-1$  or the initial hidden state at time 0, and  $r_t$ ,  $z_t$ ,  $n_t$  are the reset, update, and new gates, respectively.  $\sigma$  is the sigmoid function.

<https://pytorch.org/docs/stable/nn.html#gru>

# Exercise 12 – 2 Using GRU

## Inputs: input, h\_0

- **input** of shape  $(seq\_len, batch, input\_size)$ : tensor containing the features of the input sequence. The input can also be a packed variable length sequence. See `torch.nn.utils.rnn.pack_padded_sequence()` for details.
- **h\_0** of shape  $(num\_layers * num\_directions, batch, hidden\_size)$ : tensor containing the initial hidden state for each element in the batch. Defaults to zero if not provided.

## Outputs: output, h\_n

- **output** of shape  $(seq\_len, batch, hidden\_size * num\_directions)$ : tensor containing the output features  $h_t$  from the last layer of the GRU, for each  $t$ . If a `torch.nn.utils.rnn.PackedSequence` has been given as the input, the output will also be a packed sequence.
- **h\_n** of shape  $(num\_layers * num\_directions, batch, hidden\_size)$ : tensor containing the hidden state for  $t = seq\_len$

<https://pytorch.org/docs/stable/nn.html#gru>



# PyTorch Tutorial

## 12. Basic RNN